

The Swiss Franc Shock and Polish Voter Attitudes: Does Information Increase Support for Government Intervention?

Session 3 Analysis Report

June 3, 2026

Abstract

In January 2015 the Swiss National Bank (SNB) unexpectedly abandoned its currency floor, sending the Swiss franc sharply higher against the Polish zloty and raising monthly payments for the many Polish households holding franc-denominated mortgages. Using a nationally representative survey of 2,044 Polish adults fielded in October 2015 with an embedded randomized information treatment, we ask whether providing information about the crisis increases support for government intervention to help affected borrowers. We merge the survey with administrative demographics, estimate weighted logistic regressions of intervention support on the randomized treatments, and subject the baseline estimate to a battery of robustness checks. Factual information about the franc shock raises support for intervention by a statistically significant margin (coefficient 0.302, $p = 0.016$), while historical-context and Hungary-comparison framings do not. Direct exposure to a current foreign-currency (FX) loan is a far stronger predictor of support than any framing. The information effect is robust to weighting and to the addition of demographic controls on the full sample, but loses conventional significance once income and ideology controls shrink the analytic sample by roughly 44% — a missing-data artifact rather than confounding. We conclude that information shifts attitudes on the margin, but that personal economic exposure dominates, and that sample composition, not the controls themselves, drives the apparent fragility of the effect.

1 Introduction

The 15 January 2015 removal of the SNB’s CHF/EUR floor was a textbook external economic shock: unanticipated, large, and unevenly distributed across the population. Figure 1 shows the resulting discontinuity in the CHF/PLN exchange rate, which fed directly into the monthly payments of Polish households with franc-denominated mortgages. This report studies the political consequences of that shock using survey data described in the course codebook, drawn from Ahlquist, Copelovitch, and Walter, “The Political Consequences of External Economic Shocks,” *American Journal of Political Science* (2020).

The research question is whether *providing information* about the crisis increases support for government intervention to help affected borrowers, and whether a respondent’s own FX-loan exposure moderates that effect. Because the information framing was randomly assigned, the treatment contrasts have a causal interpretation within the surveyed population.




Figure 1: Swiss franc to Polish zloty exchange rate around the SNB’s removal of the CHF/EUR floor on 15 January 2015. The series exhibits a sharp break at the announcement before partially stabilizing over the remainder of 2015.

2 Data and Methods

2.1 Data and linkage

The analysis draws on two files. The survey itself, `swiss_franc_survey.csv`, records responses from 2,044 Polish adults interviewed between 7 and 21 October 2015, and contains the randomized information treatment, the intervention-support outcome, each respondent’s FX-loan status, a standard set of demographics, and survey weights. A companion file, `respondent_demographics.csv`, adds two attitudinal and socioeconomic measures collected separately: household income quintile and left–right ideological self-placement.

Linking the two requires care. The demographics file holds 2,994 records, of which only 1,994 correspond to actual survey respondents; the remaining thousand are identifiers that match no one, and a further 50 genuine respondents never appear in it at all. We therefore retain the survey as the master frame and attach demographics through a left join, which yields exactly the intended 2,044 rows, leaves demographics missing for the 50 unmatched respondents, and silently drops the spurious identifiers. The alternative join types each compromise the sample, either by discarding genuine respondents or by importing phantom ones. To confirm the linkage behaved as intended, we checked the final row count, verified that respondent identifiers remained unique, and reproduced the expected treatment-arm sizes (480 control, 549 information, 520 history, and 495 Hungary).

2.2 Measures and analytic sample

The dependent variable, `supports_intervention`, takes the value one for respondents who favor either “some” or “big” government intervention and zero otherwise; the 8 respondents who answered “don’t know” or left the item blank are set aside, leaving 2,036 with a usable outcome. The key independent variable is the randomly assigned framing, which took one of four forms — control, factual information, historical context, or a comparison with Hungary’s policy response — with

the control condition serving as the reference category throughout. We distinguish three levels of personal exposure to the shock: respondents currently repaying an FX loan (69), those who held one in the past (55), and those who never did (1,918). Consistent with the codebook, an income quintile of zero denotes “no income” and is retained as a substantive category rather than coded as missing. Table 1 summarizes the merged sample.

Table 1: Summary statistics, merged analytic sample.

Variable	N	Mean	SD	Min	Max
Supports intervention (0/1)	2036	0.48	0.50	0.00	1.00
Age (years)	2044	48.81	17.58	18.00	85.00
Female (0/1)	2044	0.54	0.50	0.00	1.00
Education level (1–5)	2044	2.52	1.07	1.00	4.00
Urban–rural (1–5)	2044	1.85	0.80	1.00	3.00
Income quintile (1–5)	1479	3.13	1.45	0.00	5.00
Left–right placement (0–10)	1502	0.47	1.56	-3.00	3.00
FX exposed, current (0/1)	2044	0.03	0.18	0.00	1.00
FX exposed, past (0/1)	2044	0.03	0.16	0.00	1.00

Income and ideology are missing for roughly a quarter of records by construction (Table A and the appendix). Because these two controls are the only ones with appreciable missingness, they govern how the analytic sample contracts as the specification grows — a point that proves decisive for interpreting the results.

2.3 Estimation strategy

We model the binary outcome with weighted logistic regression. Writing $p_i = \Pr(\text{supports_intervention}_i = 1)$, the baseline specification is

$$\text{logit}(p_i) = \beta_0 + \beta_1 \text{Info}_i + \beta_2 \text{History}_i + \beta_3 \text{Hungary}_i, \quad (1)$$

where each treatment indicator is measured against the omitted control group, so the coefficients recover the causal effect of each framing relative to receiving no information. We then build outward from Equation (1) in stages, adding FX-exposure indicators, the core demographic controls (age, gender, education, and urban–rural settlement), and finally the income and ideology measures. Because the model is fit by listwise deletion, we record the analytic N at every stage so that any change in the estimates can be attributed to the covariates or to the shrinking sample rather than conflated between the two. The robustness battery then perturbs the baseline along four dimensions: the presence of survey weights, the breadth of the control set, the treatment parameterization (the three arms separately versus a single “any information” indicator), and the estimation sample (all respondents versus urban dwellers only).

3 Results

3.1 Main estimates

Table 2 reports the baseline and controls-augmented weighted logistic regressions. The factual information treatment raises support for intervention relative to the control group (0.302, $p = 0.016$

in the baseline), and the estimate is stable when FX exposure and demographic controls are added on the full sample (0.324). Neither the historical-context nor the Hungary treatment produces a statistically significant effect. By far the largest coefficient belongs to *current* FX exposure (1.287, $p < 0.01$): respondents personally repaying a franc loan are dramatically more supportive of intervention, independent of any framing they received.

Table 2: Logistic regression of support for government intervention (weighted).

	(1) <i>Baseline</i>	(2) <i>+ Controls</i>
Info treatment	0.302** (0.126)	0.324** (0.127)
History treatment	0.123 (0.128)	0.133 (0.130)
Hungary treatment	0.194 (0.130)	0.185 (0.131)
FX exposed (current)		1.287*** (0.283)
FX exposed (past)		-0.163 (0.289)
Demographic controls	No	Yes
Survey weights	Yes	Yes
N	2036	2035

*** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

3.2 Robustness

Table 3 arrays six specifications side by side; the compiled version is reproduced in Figure 2. The information effect survives dropping the survey weights (column 2, 0.290, $p = 0.021$) and adding demographic controls on the full sample (column 3, 0.324, $p = 0.011$). Pooling the three framings into a single “any information” dummy yields a smaller but still significant effect (column 5, 0.218, $p = 0.042$), and restricting to urban respondents leaves the estimate essentially unchanged though borderline (column 6, 0.326, $p = 0.051$).

The one specification in which the effect loses significance is column 4, which adds the income and ideology controls. Here the coefficient remains positive and comparable in magnitude (0.267) but p rises to 0.118. The reason is sample composition rather than confounding: listwise deletion of missing income and ideology drops the analytic sample from 2,035 to 1,151 — about 44% — and the loss of precision, not a change in the point estimate, drives the result. Across every specification, current FX exposure remains large and highly significant.

Table 3: Logistic Regression: Support for Government Intervention

	(1)	(2)	(3)	(4)	(5)	(6)
	Base (weighted)	Base (unweighted)	+ Controls	Full controls	Pooled treatment	Urban only
Info treatment	0.302** (0.126)	0.290** (0.126)	0.324** (0.127)	0.267 (0.170)		0.326* (0.167)
History treatment	0.123 (0.128)	0.047 (0.127)	0.133 (0.130)	0.261 (0.174)		0.242 (0.169)
Hungary treatment	0.194 (0.130)	0.145 (0.129)	0.185 (0.131)	0.225 (0.176)		0.221 (0.169)
Any information					0.218** (0.107)	
FX exposed (current)			1.287*** (0.283)	1.251*** (0.352)	1.290*** (0.283)	1.256*** (0.319)
FX exposed (past)			-0.163 (0.289)	-0.168 (0.356)	-0.158 (0.289)	-0.049 (0.331)
Demographic controls	No	No	Yes	Yes	Yes	Yes
Survey weights	Yes	No	Yes	Yes	Yes	Yes
Sample	Full	Full	Full	Full	Full	Urban
N	2036	2036	2035	1151	2035	1208

*** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$. Reference: control group, no FX loan.

Logistic Regression: Support for Government Intervention						
	(1)	(2)	(3)	(4)	(5)	(6)
	Base (weighted)	Base (unweighted)	+ Controls	Full controls	Pooled treatment	Urban only
Info treatment	0.302 * * (0.126)	0.290 * * (0.126)	0.324 * * (0.127)	0.267 (0.170)		0.326* (0.167)
History treatment	0.123 (0.128)	0.047 (0.127)	0.133 (0.130)	0.261 (0.174)		0.242 (0.169)
Hungary treatment	0.194 (0.130)	0.145 (0.129)	0.185 (0.131)	0.225 (0.176)		0.221 (0.169)
Any information					0.218 * * (0.107)	
FX exposed (current)			1.287 * * * (0.283)	1.251 * * * (0.352)	1.290 * * * (0.283)	1.256 * * * (0.319)
FX exposed (past)			-0.163 (0.289)	-0.168 (0.356)	-0.158 (0.289)	-0.049 (0.331)
Demographic controls	No	No	Yes	Yes	Yes	Yes
Survey weights	Yes	No	Yes	Yes	Yes	Yes
Sample	Full	Full	Full	Full	Full	Urban
N	2036	2036	2035	1151	2035	1208

*** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$. Reference: control group, no FX loan.

Figure 2: Compiled robustness table (rendered from `robustness_table.tex`).

4 Conclusion

Factual information about the Swiss franc shock measurably increased Polish respondents' support for government intervention to help affected borrowers, whereas historical and comparative framings

did not. The effect is modest but robust to weighting, to control specification on the full sample, and to a pooled treatment definition. Its apparent fragility under the fullest control set is a missing-data artifact — a 44% sample reduction from listwise deletion of income and ideology — rather than evidence that the controls explain the effect away. The most substantively important finding is that personal economic exposure dominates information: respondents currently repaying a franc loan were far more supportive of intervention than any framing could induce. For applied work, the analysis is a reminder to predict and report how the analytic sample changes as controls are added, and to distinguish a vanishing point estimate from a vanishing standard error.

A Appendix: Demographic Distributions

The demographics file is summarized below over all 2,994 rows (Table A), with the full income-quintile and ideology distributions in Tables A and A and the corresponding figures following. These tables document the missingness in income and ideology that drives the sample reduction in column 4 of the robustness table.

Table 4: Respondent demographics summary (`respondent_demographics.csv`, all rows)

Statistic	Value
Total rows	2,994
<i>Income quintile</i>	
Non-missing	2,215 (74.0%)
Missing	779 (26.0%)
Modal category	Q5
<i>Left-right ideology</i>	
Non-missing	2,257 (75.4%)
Missing	737 (24.6%)
Mean	0.49
Median	0.0
Std. dev.	1.58

Table 5: Respondent income quintile distribution

Income quintile	Count	% of valid	% of all
No income	11	0.5	0.4
Q1	444	20.0	14.8
Q2	339	15.3	11.3
Q3	495	22.3	16.5
Q4	395	17.8	13.2
Q5	531	24.0	17.7
Missing	779		26.0
Total	2,994	100.0	100.0

Table 6: Respondent left–right ideology distribution

Left–right	Count	% of valid	% of all
-3	95	4.2	3.2
-2	125	5.5	4.2
-1	242	10.7	8.1
0	894	39.6	29.9
+1	303	13.4	10.1
+2	221	9.8	7.4
+3	377	16.7	12.6
Missing	737		24.6
Total	2,994	100.0	100.0




Figure 3: Distribution of household income quintile.




Figure 4: Distribution of left-right ideological self-placement.




Figure 5: Cross-tabulation of income quintile and left-right ideology.