Claude Code for Academic Researchers

A Methods Workshop in AI-Assisted Research

Brady Allardice

UPF & IBEI

About Me

Brady Allardice
PhD Candidate, Political & Social Sciences
UPF / IBEI, Barcelona
OpenAI Fellow

Research: Political economy of technological change

- Automation, AI, and political behavior
- LLM-based text classification pipelines

Background:

- M.Res. Political Science (UPF)
- M.Sc. Economics (Barcelona School of Economics)
- Former data scientist

Why this course:

- Daily user of Claude Code for academic research
- Use it across the full pipeline: data cleaning, analysis, robustness checks, code auditing, replication, writing, reviewing
- These tools can dramatically expand what a researcher can accomplish, if used with the right workflows

brady.allardice@upf.edu
github.com/bradyall

Pre-Course Survey



Please fill this out now (about 10 minutes)

- Who you are and what you research
- How you currently use AI
- What you want from the course

I use the answers to pitch the sessions and examples at the right level.

`https://docs.google.com/forms/d/e/1FAIpQLSchhltNMhTxLOL-AL-0N6_VHm8i48W6rCOn3LALyDv6dUJGkg/viewform`

Course Roadmap

Session	Focus
1	What it is and how it works: agentic coding, permissions, models, context, CLAUDE.md
2	Context strategies, parallel sessions, plan mode, and the start of verification
3	From verified data to the paper: estimation, robustness, LaTeX tables, git as the undo button, LaTeX in VS Code
4	Power use: web search, subagents, LLMs as research coders, MCPs, skills, and hooks

4 sessions $\times$ 2 hours | about half hands-on | Pair programming throughout

How We'll Work: Your Own Workspace

The course materials live in one shared repo, organized as self-contained session folders:

```
claude-code-workshop/
  session_1/ <- today's materials live here
  session_2/
  session_3/
  ...
```

The workflow, every session: the start-of-session command copies that session's folder out of the repo into your own workspace, and you work there. The shared repo stays read-only for you.

Why

I can push new material anytime without breaking your work. You can edit, break, delete, and redo without worrying about the shared state. No merge conflicts.

The Start-of-Session Ritual: One Command

Do this at the start of every session. Open **github.com/bradyallardice/claude-code-workshop** and copy the command at the top of the page.

Mac (Terminal)

```
curl -fsSL https://raw.githubusercontent.com/
  bradyallardice/claude-code-workshop/student/
  tools/start_session.sh | bash -s -- 1
```

Change the last number to today's session.

Windows (PowerShell)

```
irm https://raw.githubusercontent.com/
  bradyallardice/claude-code-workshop/student/
  tools/start_session.ps1 | iex
```

It asks which session.

- Updates the course materials to **today's version**, backing up anything you changed
- Copies today's folders into `my_workspace`. On later runs it **never overwrites a file you edited**: the newer version appears beside it as `<name>.new`
- **Work only in** `my_workspace`, never inside the course repo (it is reset each time)
- Opens `my_workspace` in VS Code: **start Claude Code there**. This is the folder every exercise runs from.
- **Lost it?** `my_workspace` is in your home folder (the READY box prints the path). Mac: Finder → Go → Home. Windows: `%USERPROFILE%` in the File Explorer address bar. Or VS Code → File → Open Recent
- **Check:** the *Materials version* in your READY box matches mine

If anything fails: `tools/TROUBLESHOOTING.md`.

Setup Check

Open a terminal inside VS Code, in the window the start command opened (`my_workspace`): menu **Terminal** → **New Terminal**, or press `Ctrl+`` . A panel opens at the bottom of the window. Type each command there and press Enter:

1. Verify Claude Code: `claude --version`
2. Launch Claude Code: `claude`
   Try a simple message: `What files are in this directory?`
3. Compile the test document: open `setup_check.tex` and click the green **Build** arrow (top right)
   A PDF with a small table should open beside it. Terminal alternative:
   `latexmk -pdf setup_check.tex`

If you completed pre-work: help your neighbor.
If you're stuck: pair with someone who is set up and keep going. Tell me, and I'll come round during the first exercise.
LaTeX problems must be fixed before Session 3, which has no time to absorb them.

Core Framework

By the end of this course, you will be able to do three things:

1. Integrate

Use Claude Code (or Codex, or whichever coding agent comes next) inside a **real academic workflow**: data, code, analysis, writing, and version control in one place.

2. Develop core AI skills

Learn what determines the quality of what you get out: **how you instruct it, the context it has, and the model you choose**.

3. Stay in charge

Do it in a way that keeps **you in control** and keeps your work **accurate and verifiable**.

The question is never "can the agent do this." It almost always can.

The question is: **can you get it done productively, and can you show it is right?**

You are the principal investigator and the agent is the research assistant. You direct the work, make the analytical decisions, and verify the results.

Part A

What This Is and What You Can Do With It

Demo + Exercise

Live Demo: The Task

Goal: Link the election and demographic data, run a regression, and write up the results in a LaTeX document.

This requires:

1. Reshape election data to one row per county-year with vote shares
2. Aggregate 15M person-level records to county-level means using survey weights
3. Harmonize FIPS codes across datasets (different formats)
4. Merge on county $\times$ year and validate the join
5. Run a regression relating election outcomes to county demographics
6. Produce the regression output (a table)
7. Summarize the results in a LaTeX document

In Stata or R, this takes well over an hour of careful coding.

Live Demo: The Data

Dataset 1: County Presidential Election Returns

- Source: MIT Election Data + Science Lab
- County-level vote totals by candidate and party, 2000–2024
- $\sim$94,000 rows, one row per county $\times$ year $\times$ candidate
- Variables: FIPS codes, candidate names, party, vote counts, vote mode

Dataset 2: IPUMS American Community Survey Microdata

- Source: IPUMS USA (University of Minnesota)
- Person-level survey records with demographic, economic, and geographic variables
- $\sim$15 million rows per year, one row per person
- Variables: age, education, employment, income, race, poverty status, county FIPS

Live Demo: Same Task, Two Ways

We will do this task twice, slowly, so you can see exactly what changes.

Part 1: In a chat (Claude or ChatGPT)

- Same task, in the web chat
- What does it give back?
- What are you left to do by hand?

Part 2: In Claude Code

- First: **where the data lives** on disk
- Then: **the exact prompt** I type
- Then: watch Claude Code work

Watch for: where did the data go, who ran the code, and how would you check the answer?

What Just Happened

It wrote the code:

- Read both files
- Identified the join key
- Handled cleaning steps
- Wrote a complete script
- Saved the new code to a file

It also:

- Understood the data structure (person-level vs. county-level, survey weights, FIPS formats)
- Ran the code and diagnosed errors
- Provided summaries of the output (row counts, variable descriptions, merge diagnostics)

But it also made a decision about our data without asking.

Both of these things are true:

1. It took 90 seconds instead of 45 minutes
2. It made analytical choices autonomously

The Coding Paradigm Spectrum: Manual Coding

Manual Coding

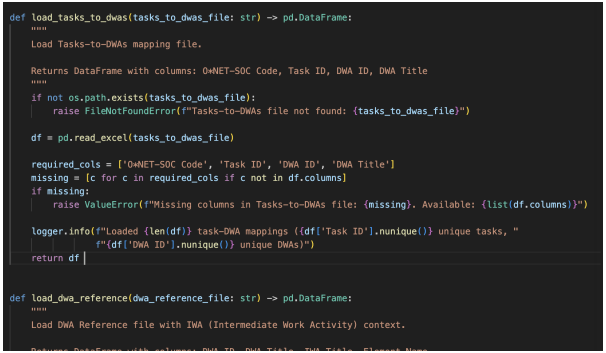
Autocomplete / Copilot

Web Chat

Agentic Coding

You write every line. Full control, no assistance.

This is what you do now in Stata or R.



```python
def load_tasks_to_dwas(tasks_to_dwas_file: str) -> pd.DataFrame:
    """
    Load Tasks-to-DWAs mapping file.

    Returns DataFrame with columns: O*NET-SOC Code, Task ID, DWA ID, DWA Title
    """
    if not os.path.exists(tasks_to_dwas_file):
        raise FileNotFoundError(f"Tasks-to-DWAs file not found: {tasks_to_dwas_file}")

    df = pd.read_excel(tasks_to_dwas_file)

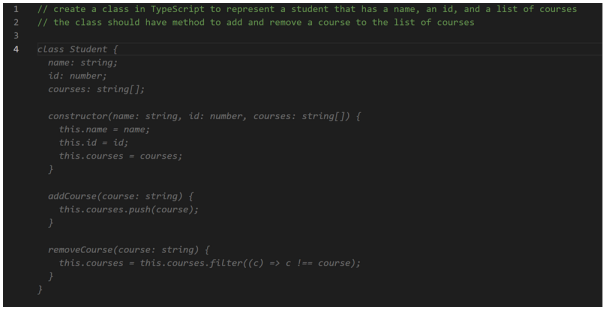
    required_cols = ['O*NET-SOC Code', 'Task ID', 'DWA ID', 'DWA Title']
    missing = [c for c in required_cols if c not in df.columns]
    if missing:
        raise ValueError(f"Missing columns in Tasks-to-DWAs file: {missing}. Available: {list(df.columns)}")

    logger.info(f"Loaded {len(df)} task-DWA mappings ({df['Task ID'].nunique()} unique tasks, "
                f"{df['DWA ID'].nunique()} unique DWAs)")
    return df


def load_dwa_reference(dwa_reference_file: str) -> pd.DataFrame:
    """
    Load DWA Reference file with IWA (Intermediate Work Activity) context.
```

The Coding Paradigm Spectrum: Autocomplete / Copilot

- Manual Coding
- **Autocomplete / Copilot**
- Web Chat
- Agentic Coding



```
// create a class in TypeScript to represent a student that has a name, an id, and a list of courses
// the class should have method to add and remove a course to the list of courses

class Student {
  name: string;
  id: number;
  courses: string[];

  constructor(name: string, id: number, courses: string[]) {
    this.name = name;
    this.id = id;
    this.courses = courses;
  }

  addCourse(course: string) {
    this.courses.push(course);
  }

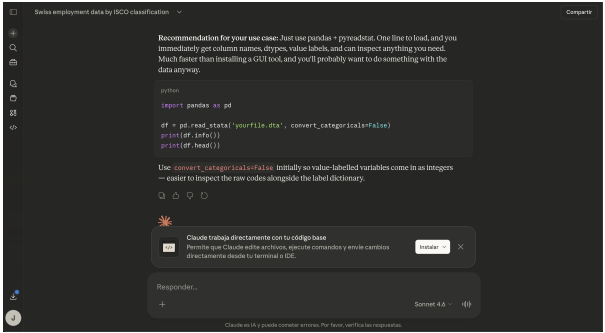
  removeCourse(course: string) {
    this.courses = this.courses.filter((c) => c !== course);
  }
}
```

AI suggests the next line as you type.

You accept or reject. The AI assists; you drive.

The Coding Paradigm Spectrum: Web Chat

- Manual Coding
- Autocomplete / Copilot
- **Web Chat**
- Agentic Coding



You describe a task; AI gives you code to copy-paste.

The AI cannot see your project.

The Coding Paradigm Spectrum: Agentic Coding

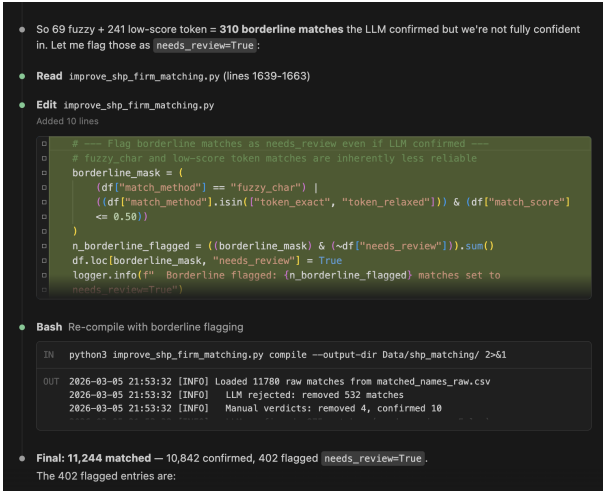
Manual Coding

Autocomplete / Copilot

Web Chat

Agentic Coding

AI reads, writes, and runs code **in your project**.



Why Agentic Coding Is Different

Paradigm	Pros	Cons
Manual Coding	Full control	Slow, everything by hand
Autocomplete / Copilot	You control logic and structure	Still line-by-line
Web Chat	Can handle complex tasks	Copy-paste; AI misses context
Agentic Coding	**Fastest; works on your project**	**Can act autonomously**

The Key Insight

Agentic coding is the first paradigm where the tool **interacts with your project directly and autonomously**. That is what makes it so productive, and why learning to use it well matters.

These principles apply to Claude Code, Cursor, Codex, Windsurf, or whatever comes next.

Claude Code vs. Cowork vs. Web Chat

Claude Code and Cowork are **agentic**: they act on your files. A web chat with code execution on works on uploaded copies.

	Claude Code	**Claude Cowork**	**Web chat**
Built for	Writing, running, debugging code	Documents, spreadsheets, files	Conversation; analysis of uploaded files
Interface	Terminal, VS Code, desktop app, web	Claude desktop app	Browser or app
Your files	Stay in your project folder	Folders you share with the app	Uploaded copies (size limits)
Output	Scripts and a project you can re-run	Finished files (reports, slides)	An answer, with code, charts or files
What you see	Every command and edit	Task progress	The reply; code may be hidden

Why this course uses Claude Code

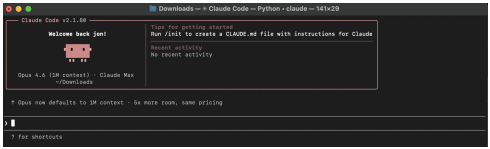
It leaves behind **scripts you can inspect, re-run, and version**, in your own project, so every step from raw data to table is auditable. A finished file or a chat answer does not, by default.

Where Does Claude Code Live?

Claude Code runs **inside your terminal**, not in a browser, not in an app.

1. You open a terminal
2. Navigate to your project: `cd my_project/`
3. Type `claude` and press Enter
4. You're now talking to Claude Code

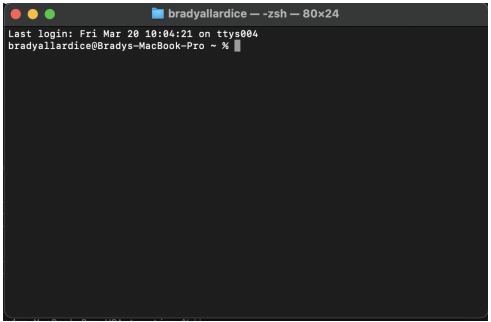
It can see your files, read your data, write scripts, and run them, all without copying and pasting.



This is what Claude Code looks like when you launch it.

What Is the Terminal?

The terminal is a text-based interface to your computer. Instead of clicking, you type commands.



Claude Code writing and running code in the terminal.

Essential commands:

```
# Where am I?
pwd

# What's in this folder?
ls

# Move into a folder
cd my_project/

# Run a Python script
python my_script.py

# Launch Claude Code
claude
```

Same Claude Code, Different Front Doors

Claude Code is **one agent**. You can reach it from several places, and it behaves the same in all of them:

Where	What it is / when to use it
Terminal (CLI)	The original. Run `claude` in any project folder — full agent, always available, great for quick jobs or tools without an editor plug-in.
VS Code / JetBrains	**Our home base for this course.** The same agent inside your editor: inline diffs, click-to-open file references, your code and Claude side by side.
Desktop app	Claude Code in its own window (Mac/Windows) instead of a terminal.
Web (claude.ai/code)	Claude Code in the browser for cloud sessions — nothing to install.

Not the same as the Claude app

The **Claude app** (claude.ai or the desktop chat) is a conversational assistant: you paste text in and copy answers out. It can't see your project, run your code, or edit files in place. **Claude Code** acts directly in your repo — reading files, editing them, and running commands for you.

Terminal vs. VS Code

The terminal is the original way to use Claude Code, and it still works anywhere.

But for this course, we'll be working in **VS Code**: the file you're editing, the code Claude is writing, and the conversation all live in one window. This holds for **Python, R, and Stata** alike, since VS Code has extensions for all three (next slides).

Why VS Code?

VS Code is a free code editor with a built-in terminal.

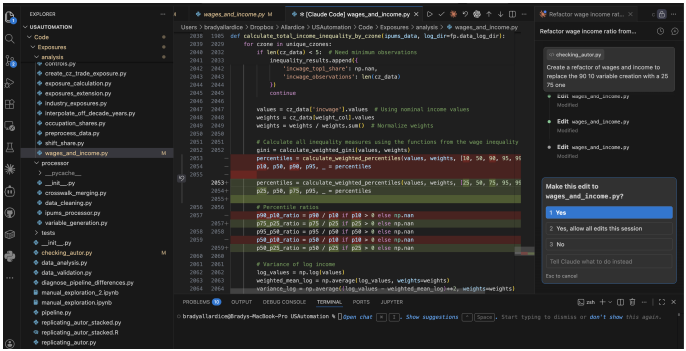
Everything in one window:

- **Left:** file explorer
- **Center:** your code
- **Right:** Claude Code

You see what Claude Code writes *as it writes it*.

Today we use three features:

1. Open a folder
2. View files in the editor
3. Talk to Claude Code

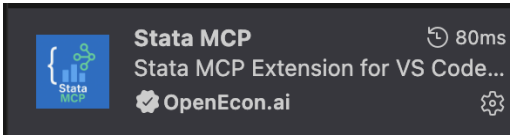


Coming from RStudio? ▸ Configure VS Code to look like RStudio

VS Code Runs Your Code Too

VS Code isn't just an editor — with the right extensions, it executes your scripts directly:

- **Python** (`.py`, `.ipynb`) — run line-by-line with Shift+Enter, or run the whole file
- **R** (`.R`, `.Rmd`) — same workflow as RStudio, same interactive console
- **Stata** (`.do`) — run do-files via the Stata extension, results in the integrated terminal

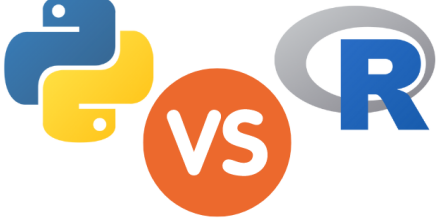


The Stata extension for VS Code.

One editor, one window: Claude writes the code, and you run it without leaving the screen.

Why Python

- Claude Code is substantially better at Python than Stata or R
- The principles (task decomposition, verification, control) transfer to any language
- Python is increasingly where the tools and ecosystem are for computational research



Python Files: `.py` vs. `.ipynb`

Notebooks (`.ipynb`)

```
df_ipums_pre_exposure = pd.read_pickle(f'/Users/bradyallardice/Desktop/PhD/Projects/USAutomation/Data/Merge/IPUMS/processed_ipums_1040.pkl')
df_ipums_pre_exposure.groupby(['occ1950rj','czone'])['perwt_by_afacter'].sum().reset_index()
```

[29]

Python

	occ1950rj	czone	perwt_by_afactor
0	2.0	3700.0	62.676506
1	2.0	3800.0	37.315493
2	2.0	4200.0	9.041921
3	2.0	4301.0	12.407893
4	2.0	4702.0	12.759623
...	...	...	...
68909	970.0	39206.0	238.110368
68910	970.0	39301.0	83.274260
68911	970.0	39302.0	1712.524776
68912	970.0	39303.0	1064.010787
68913	970.0	39400.0	25848.531526

68914 rows × 3 columns

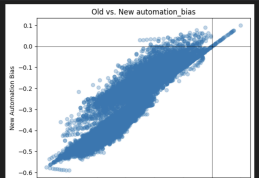
Visualizing Automation Exposure Changes

```
import matplotlib.pyplot as plt

fig, ax = plt.subplots()
ax.scatter(df_final['automation_bias_per_worker'], df_final['automation_bias_per_worker_ffilled'], alpha=0.3)
ax.set_xlabel('Old Automation Bias')
ax.set_ylabel('New Automation Bias')
ax.axhline(0,color='k',linewidth=.5); ax.axvline(0,color='k',linewidth=.5)
ax.set_title('Old vs. New automation_bias')
plt.show()
```

[30]

Python



Scripts (`.py`)

Code > Exposures > checking_autor.py > ...

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import pearsonr, spearmanr
import os
import pickle
from pathlib import Path

# Set up plotting style
plt.style.use('ggplot')
sns.set_context("talk")

import General.filepaths as fp

def load_data_files():
    """Load all intermediate data files needed for diagnostics"""
    data_files = {}

    # Try to load cz-level data
    try:
        occ_exposure_path = f'{fp.exposures_out_dir}occ_ind_4080rj_panel_cw_t85.dta'
        data_files['occ_exposure'] = pd.read_stata(occ_exposure_path)
        data_files['occ_exposure'] = data_files['occ_exposure'][['year','occ1950rj','N
        print(f"✓ Loaded occupation exposure data from {occ_exposure_path}")
        print(f"  Shape: {data_files['occ_exposure'].shape}")
    except FileNotFoundError:
        print(f"✗ Could not find occupation exposure data at {occ_exposure_path}")

    # Try to load occupation shares data
    try:
        occ_shares_path = f'{fp.exposures_out_dir}final_occ_shares_all_years_by_cz.csv
        data_files['occ_shares'] = pd.read_csv(occ_shares_path)
        print(f"✓ Loaded occupation shares data from {occ_shares_path}")
        print(f"  Shape: {data_files['occ_shares'].shape}")
    except FileNotFoundError:
        print(f"✗ Could not find occupation shares data at {occ_shares_path}")
```

Why Scripts for This Course

Notebooks (`.ipynb`)

- Cells you run one at a time
- Great for exploration and teaching
- State depends on which cells you ran and in what order
- Hard for Claude Code to work with because it cannot "click" cells

Scripts (`.py`)

- Plain text file, runs top to bottom
- Reproducible: `python my_script.py` always does the same thing
- Claude Code can read, write, and run them directly
- **This is what we use today**

The Key Difference

Claude Code operates on files, not interactive sessions. Scripts run end-to-end and produce the same result every time, exactly what you need for reproducible research.

What Changes When You Use Scripts

You can run scripts line by line in VS Code (Shift+Enter), just like a notebook. But scripts aren't really built for that.

Their strength is that they run **end-to-end**: `python my_script.py` always does the same thing, top to bottom. That is what makes them reproducible, and what Claude Code needs.

What this means for your workflow:

- More **checkpointing**: save intermediate datasets to disk so you can inspect them
- More **logging**: print row counts, shapes, and summaries as the script runs
- More **intermediate file saves**: write CSVs between pipeline stages
- Overall **less interactive**, more structured

This feels slower at first. But it produces pipelines that are reproducible, debuggable, and that Claude Code can work with directly.

Virtual Environments

What is a virtual environment?
A self-contained folder that holds a separate copy of Python and its packages. Each project gets its own environment, so packages from one project don't interfere with another.

Why you need one:

- Without it, every project shares the same packages, and updating one project can break another
- Claude Code installs packages when it needs them. You want that isolated, not touching your system Python
- Reproducibility: you know exactly which packages your project uses

Think of it like this: a virtual environment is a clean desk for each project. You only put the tools you need on it, and cleaning up one desk does not affect the others.

Setting Up a Virtual Environment

Run these commands in your terminal, from your project folder:

```
# 1. Create the environment (once per project)
python3 -m venv AIAgents

# 2. Activate it (every time you open a new terminal)
source AIAgents/bin/activate      # Mac/Linux
AIAgents\Scripts\activate         # Windows

# 3. Install packages
pip install pandas numpy statsmodels

# 4. Confirm it's active (you should see (AIAgents) in your prompt)
which python
```

In VS Code: once you create the `AIAgents` folder, VS Code will detect it and ask to use it as your interpreter. Say yes. After that, every terminal it opens will activate the environment automatically.

What Claude Code Can Do for Academic Research

Data Work

- Collect data: scrape websites, call APIs, download and organize sources
- Clean, reshape, and merge datasets
- Construct variables from raw data
- Build panel datasets across sources
- Handle missing data, recoding, harmonization

Analysis

- Exploratory data analysis and summary statistics
- Generate robustness checks and specification variants
- Create publication-quality visualizations
- Synthesize results across specifications

Writing and Review

- Draft methods sections, data descriptions, results summaries
- Audit existing scripts for silent errors
- Review code against your methods section
- Write documentation, READMEs, replication packages

Research Infrastructure

- Process text corpora at scale (NLP, classification)
- Crosswalk and linkage construction
- Refactor and optimize slow pipelines

The RA Analogy

How to Think About Claude Code

Claude Code is like a research assistant who is:

- Extremely fast
- Has read a lot
- Will **never** say "I don't know how to do that"

Instead, it will just do it wrong and hand it in.

This course teaches you to capture the value while managing the risk.

A Tool Built for Software Engineers

Claude Code was designed for professional developers. Their world is different from ours.

Software Engineering

- "Works" = **runs correctly**. Clear success criteria: tests pass, feature ships, users are happy
- Code is the product
- Errors are caught by automated tests
- Speed and iteration are paramount

Academic Research

- "Works" = **???**. Code that runs without errors can produce completely wrong results. Output that looks plausible can be based on the wrong sample, wrong specification, or wrong variable
- Code is the *instrument*; the paper is the product
- Errors are caught by *you*, or by Reviewer 2
- Correctness matters more than speed

The defaults, instincts, and shortcuts built into these tools reflect engineering priorities.
Using them well for research means overriding some of those defaults.

Do You Own It?

Pick any result your code produced. Without looking at the code:
what got dropped, what got merged, and why?

If you can't answer that in 30 seconds, you don't own it.

Before you delegate:

- What would wrong look like? Can I describe a specific failure mode before I run this?

After you delegate:

- Can I describe what this code does without reading it?
- Can I state what decisions shaped the output and why?

You didn't write it. That's fine.
But every decision in it is yours. **Can you justify it as if you chose it yourself?**

Why Verification Matters
What happens when you skip verification:

1. **Silent observation loss:** A merge silently drops 300 observations.
   The code looks clean. The error is invisible without checking row counts.
2. **Invalid transformation:** A log transformation applied to a variable containing zeros.
   No error raised, wrong results.
3. **Incorrect panel operations:** A lagged variable computed without sorting by time within panels.

These are the kinds of errors that end up in published work.

The Control Principle

You Are the PI. Claude Code Is the RA.

The question is not "can Claude Code do this" but "what is the best way to get this done reliably."

If you cannot verify a result, you should not delegate the task.

What's Ahead

By the end of this course, you will have:

- Built a data panel in plan mode and verified it against expectations you wrote first *(Session 2)*
- Estimated a model, run robustness variants, and turned the results into a LaTeX table and paper section, with git as your undo button *(Session 3)*
- Used web search and subagents to scan literature and code text, and checked every citation and code *(Session 4)*
- Connected tools such as Zotero and built your own skill *(Session 4)*

You will know which parts Claude Code should handle and which require your judgment.

Now that you have seen what Claude Code can do,
let us look at how it actually works.

The model, the context, and the skills that determine what you get out of them.

How Claude Code Works: Step by Step

Every request runs through the same sequence.

Once, before you type anything

1 **Load** the standing instructions and settings it starts with, and the tools it may use

The loop: repeats until the task is done

2 **Read** your prompt
3 **Gather** open files and search the project for what is relevant
4 **Reason** decide the next step
5 **Act** edit a file or run a command, asking your permission where required
6 **Observe** read the output or error that comes back, then return to step 3

When it stops

7 **Report** what it did and what it found

Everything lands in context

Each file read, command run, and result is added to the context window, which is finite.

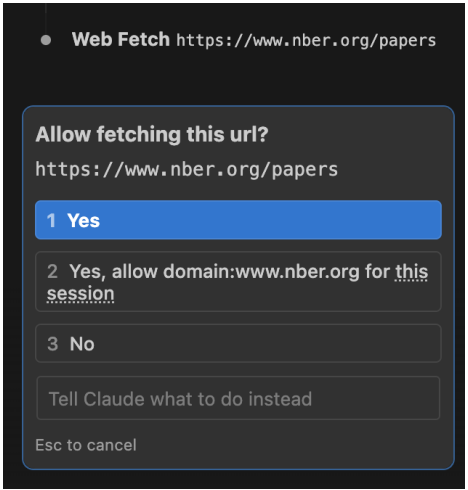
Opening Claude Code Inside the Project

With your project folder open in VS Code, there are four ways to open Claude Code. They all open the same agent:

Where	Detail
Editor toolbar	Spark icon, top-right of the editor. Only appears when a file is open.
Activity Bar	Spark icon in the left sidebar. Opens your list of sessions.
Status Bar	"Claude Code" at the bottom right. Works even with no file open.
Command Palette	`Cmd/Ctrl+Shift+P`, then type "Claude Code".

If you cannot find the icon, you probably have no file open. Use the Status Bar or the Command Palette.

When Claude Code Asks for Permission



Before it acts (step 5 of the loop: running code, writing files), Claude Code asks you.

Your three options:

1. **Yes**: allow this one time
2. **Yes, always allow**: allow this type of action for the rest of the session (no more prompts)
3. **No**: reject, and tell Claude Code what to do differently

"No" is not a dead end.

Use it to redirect: *"No, use a left join instead"* or *"No, don't overwrite that file, create a new one."*

Exercise: Your First Commands

Pick your language. The starter script builds a county-year vote-share panel — find it in `session_1/scripts/` (open `my_workspace` in VS Code and run it from there, not from inside `session_1`):

- **Python:** `exercise_build_panel.py`
- **R:** `exercise_build_panel.R`
- **Stata:** `exercise_build_panel.do`

Task 1: Ask Claude Code to explain what your script does, line by line.

Task 2: Ask it to add verification checks, documentation, and print statements.

Task 3: Ask it to change the two-party vote shares to total vote shares.

Exercise Debrief

- What did it get right?
- Anything wrong or confusing?
- That is normal and expected

Common Actions and Risk Levels

Action	Example you'll see	What it does	Risk level
`Read`	`Read scripts/clean.py`	Looks at a file	Safe
`Glob`	`Glob data/**/*.csv`	Finds files by name	Safe
`Grep`	`Grep "county_fips"`	Searches inside files	Safe
`Edit`	`Edit clean.py line 42`	Changes existing code	Review first
`Write`	`Write scripts/merge.py`	Creates a new file	Review first
`Bash`	`python scripts/clean.py`	Runs a command	Read before approving
`Bash`	`pip install pandas`	Installs a package	Read before approving
`Bash`	`rm data/old_panel.csv`	Deletes a file	Read before approving
`Bash`	`mv output.csv data/`	Moves a file	Read before approving
`Bash`	`git push`	Uploads to remote server	Read before approving

Safe to "always allow":

Reading and searching. Claude Code is just looking, not changing anything.

Always read before approving:

Bash commands can do *anything*: run code, delete files, install software. Check what it is actually running.

Restricting Claude to This Project

File access is controlled via `.claude/settings.json` — already configured in this course:

```
{
  "permissions": {
    "allow": ["Read(/**)", "Edit(/**)", "Read(~/.claude/**)"]
  }
}
```

What this does:

- Allows reads and edits *within* the project directory.
- Allows Claude to load your user-level config (`~/.claude/`).
- Anything outside the project *prompts for permission* rather than being silently allowed.

Protecting Sensitive Files

"Always allow reads" assumes there is nothing sensitive to read.
If your project has restricted data or IRB-protected files, you need to keep them away from Claude Code.

For now, the reliable rule is simple:

- Don't keep restricted or identifiable data inside the folder you run Claude Code in.
- Point Claude at a working copy with the sensitive files or columns removed, or keep that data in a separate location entirely.
- When in doubt, don't let it read the file — review before you approve.

Coming in Session 4

There are dedicated guardrails — ignore rules and permission denies — for hiding files from Claude. We'll set them up properly, and see where the naive version quietly fails, once you have a real project.

Getting Files In: Drag and Drop

You don't have to describe a file in words — you can hand it straight to Claude.

Hold `Shift` and drop into the chat panel (in VS Code) to add something as context:

- A **data file** (`.csv`, `.xlsx`) so Claude can see the columns
- A **figure or PDF** — a paper, a table, a chart to reproduce
- A **screenshot** of an error message or a plot that looks wrong

Two ways to do it

- **Shift+drag** a file from the VS Code Explorer (or your Finder/desktop) onto the chat box — holding `Shift` is what drops it in as context instead of opening it
- **Paste** an image directly with `Cmd+V` / `Ctrl+V` — great for screenshots

A screenshot of a broken plot plus "why does this look wrong?" is often faster than describing it.

AI Skill Gap

Using AI is not using it well: 758 BCG consultants, randomly assigned

	Inside AI's frontier: quality	Outside: share correct
No AI	baseline	84%
GPT-4 only	+38%	70%
GPT-4 + prompt-engineering overview	+42.5%	60%

The same pattern elsewhere

- **Cognizant survey (self-reported):** 64% of trained vs. 36% of untrained report 20%+ productivity gains
- **PwC job postings:** AI-skill wage premium rose from 14% (2024) to 62% (2026)

The takeaway

Same tool for everyone. Training raised quality where AI was reliable and lowered accuracy where it was not. Knowing the difference is the skill.

Only HBS/BCG is experimental (Dell'Acqua et al. 2023). Overview vs. GPT only: $p = .03$ (quality), .08 (accuracy).

The AI Skill Gap

Two people can use the same tool on the same task and get very different results. **Output quality is not a property of the model alone.**

1. The model

- Which model you run
- How much it is allowed to think before answering
- Its limits and failure modes

2. The context

- Claude Code sees your repo; web chat only what you paste
- What it is told about your project (`CLAUDE.md`)
- What has built up in the conversation

3. Your skill

- How precisely you specify the task
- How you break a problem into steps
- Whether you can **tell when the output is wrong**

The takeaway

A capable model does not guarantee a good result. Two of the three factors are **yours**, and the third is a choice you make. The tool will not close this gap for you; practice does.

Choosing a Model

Claude Code lets you switch between model tiers.
Each trades off **quality**, **speed**, and **cost**.

	Haiku	**Sonnet**	**Opus**	**Fable**
Speed	Fastest	Fast	Slower	Slowest
Cost	Lowest	Moderate	High	Highest
Reasoning	Basic	Strong	Very strong	Strongest

Fable is the top tier, above Opus, where it is available to you.

To switch models in Claude Code:
Type `/model` and select, or press `Option+P` (Mac) / `Alt+P` (Windows, Linux).

When to Use Each Model

Haiku

Fast, cheap, good enough for simple tasks.

- Reformatting files
- Renaming variables
- Generating boilerplate
- Quick questions about syntax

Sonnet

The everyday workhorse for most research tasks.

- Writing data cleaning scripts
- Building panels and merges
- Debugging code
- Drafting documentation

Opus

Strong reasoning; use when it matters.

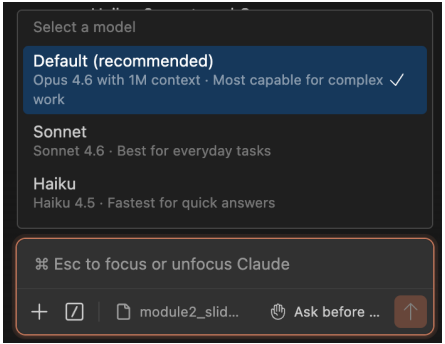
- Complex multi-file refactors
- Subtle specification decisions
- Reviewing analytical code
- Large-context projects

Fable

Top tier; for the hardest problems.

- Tasks where Opus still gets it wrong
- The most demanding reasoning and review
- When cost is not the constraint

Model Choice: The Research Perspective



The Key Tradeoff

Smarter models are slower and more expensive.
A wrong answer fast is worse than a right answer slow.

Rules of thumb:

- **Default to Sonnet**: best balance of quality and speed
- **Upgrade to Opus** for careful reasoning or many files
- **Fable** for the hardest problems, if you have it
- **Drop to Haiku** for mechanical, easy-to-verify tasks

Effort: How Hard It Works

Beyond *which* model, you control *how hard* it works. **Effort** is the dial: it scales how much Claude reasons, reads, runs, and writes.

Lower effort

Reasons little, makes fewer tool calls, answers tersely. Fast and cheap.

- Reformatting, renaming
- Simple, well-specified edits
- Quick factual questions

Higher effort

Reasons more, reads and checks more *before* acting, explains in more detail. Slower, more tokens, fewer mistakes.

- Debugging stubborn errors
- Specification decisions
- Planning a multi-step change

When in doubt

If the task involves a **judgment call** or Claude got it wrong once already, raise the effort. For mechanical work, keep it low.

Dialing the Effort

The effort setting. Set it with `/effort` (a slider) or from the slider in `/model`. From least to most thorough:

`low` → `medium` → `high` → `xhigh` → `max`

Higher effort = more reasoning and tool calls: thorough, slower, costlier. It is set once and stays for the session. Not every model offers `xhigh`, and `/effort auto` clears your choice.

One-off boost. Put the word `ultrathink` in a single prompt to get deeper reasoning for *that turn only*. Your session effort does not change. Older cues like "think" and "think hard" are no longer special; Claude reads them as ordinary words.

`ultracode` **is not a level.** It is a separate toggle (`/effort ultracode`) that makes Claude orchestrate many agents on one task. It can consume a large amount of tokens, and you will not need it in this course.

Effort is not free. It spends tokens and time — reach for it when correctness matters, not by reflex.

Where Effort Acts in the Loop

Once, before you type anything

1 **Load** the standing instructions and settings it starts with, and the tools it may use

The loop: repeats until the task is done

2 **Read** your prompt

3 **Gather** open files and search the project for what is relevant

4 **Reason** decide the next step: **the effort setting controls how hard**

5 **Act** edit a file or run a command, asking your permission where required

6 **Observe** read the output or error that comes back, then return to step 3

When it stops

7 **Report** what it did and what it found

Effort acts across the loop, not only at step 4

It scales how much Claude reasons (4), how many files it reads and commands it runs (3, 5, 6), and how detailed its report is (7). Lower effort means fewer tool calls and terser answers.

Cost Management

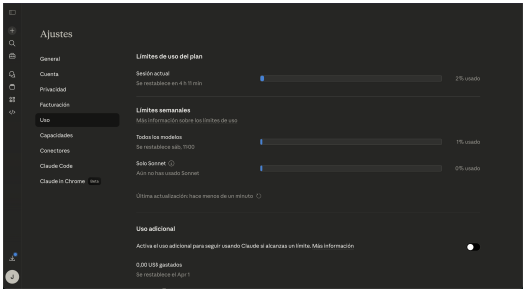
A Real Constraint

On the \$20/month plan, you can burn through your allocation in 30–45 minutes of active use if you are not strategic.

Principles:

1. Not everything needs Claude Code. Small edits are faster by hand
2. A well-written CLAUDE.md reduces back-and-forth (= fewer tokens)
3. If Claude Code is spinning on a wrong approach, **intervene early**
4. Batch related tasks into one well-scoped request
5. Plan mode adds tokens but prevents costly wrong approaches

Usage Limits and Model Cost



How limits work:

- Token allowance resets on a **5-hour rolling window**
- A separate **weekly cap** limits total usage
- When you hit either limit, Claude Code pauses until the window resets

Which models burn tokens fastest:

- **Opus** uses $\sim 5\times$ more tokens per task than Haiku
- **Sonnet** is the middle ground
- **Haiku** stretches your budget the furthest

Match the model to the task.

Don't use Opus to rename variables.

What Is Context?

Context is everything Claude can see when it writes a response. It works from nothing else.

What is in it

- Your instructions and the files it reads
- Standing instructions loaded at the start
- The conversation so far: its earlier answers and the code it ran

How you use it

- State the task, the inputs, and the constraints
- Point it to the right files, codebooks, and documentation
- Write conventions down once (`CLAUDE.md`)
- Keep tasks focused

What is not in context does not exist for Claude

It does not know your data, your field's conventions, or what you decided last week, unless that is written somewhere it can read.

More Context, Better Execution

The more relevant context Claude has, the better it executes the task. That is why running it on your computer, with direct access to your files, is so valuable.

Web chat

- Sees only what you paste or upload
- You decide what is relevant and copy it in
- Guesses at the rest: column names, file layout, what the existing script does

Claude Code on your computer

- Reads your data, scripts, codebooks, and `CLAUDE.md` directly
- Runs the code and sees the real errors and output
- Checks its work against the actual files

The advantage is not primarily a smarter model. It is **access to your project**.

More means *relevant*. The context window is finite, which is the next slide.

The Context Window

Claude Code can only "see" a limited amount of text at once.

The context window includes:

- Your instructions
- The files it reads
- Its own responses

When a conversation gets long or a project gets large, it loses earlier context.

Why the Context Window Matters

The Consistency Problem

In a complex project with many files, Claude Code may change one file in a way that is **inconsistent** with another file it can no longer see.

Errors compound when the agent loses track of earlier decisions.

We will cover strategies for managing this in Session 2.

Practical Implication

Think in Projects, Not Sessions

You need to think in terms of:

- Project directories
- File paths
- Scripts that run end-to-end

Not interactive exploration.

Where Claude Looks for Context

Claude Code reads instructions from **several levels** when it starts a session. Anything it finds is loaded into the context window before your first message.

- **User level** — `~/.claude/CLAUDE.md`
  Applies to *every* project on your machine. Personal preferences, defaults, things you want Claude to know about you regardless of what you are working on.
- **Project level** — `CLAUDE.md` in the project root
  Checked into git, shared with collaborators. Project conventions, data structure, constraints.
- **Subfolder level** — `CLAUDE.md` inside a subdirectory
  Loaded when Claude works in that subtree. Useful for scoped rules (e.g. a `cleaning/` folder with its own conventions).

How Claude Reads the Hierarchy

When Claude Code starts, it walks from the user level down to your working directory and loads **every applicable CLAUDE.md** it finds. They are concatenated into context, with more specific (deeper) files appearing later — so they take precedence when guidance conflicts.

This hierarchy is everywhere

The same user → project → subfolder pattern shows up across the power-user sessions:

- **Skills** — defined at the user or project level, loaded on demand
- **Hooks** — configured per user or per project
- **MCP servers & permissions** — scoped per user or per project
- **Subagents** — registered at the user or project level

Learn the hierarchy once and it tells you where every other piece of configuration belongs.

Everything Claude Reads Before You Type

Loaded in full at startup

- `CLAUDE.md`: user, project root, and parent folders
- `CLAUDE.local.md`: private notes
- **Imports**: files pulled in with `@file`
- `.claude/rules/` files (no `paths:` filter)
- **Memory**: first 200 lines of `MEMORY.md`

Only a name and description

- **Skills**: description only; full text loads when used
- **Subagents**: name and description
- **MCP servers**: tool names only

On demand, when relevant

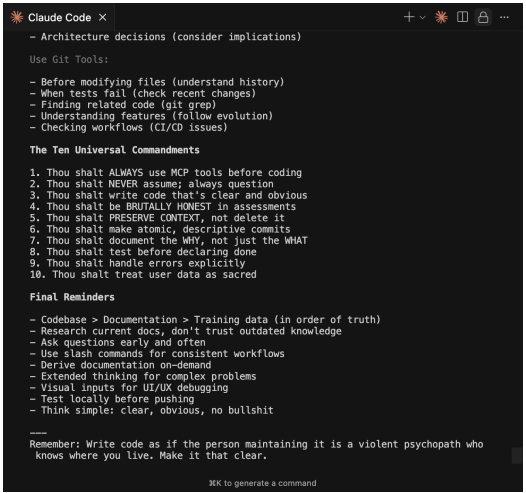
- Subfolder `CLAUDE.md`: when Claude works in that folder
- Rules with a `paths:` filter
- Memory topic files
- Full skill text and MCP tool definitions

Also always there: Claude's built-in system prompt, your output style, and a snapshot of the working directory and git status.

See it for yourself

Type `/context` to see what is loaded and how much room it takes. Type `/memory` to open the memory and CLAUDE.md files.

CLAUDE.md: Standing Instructions for the Agent



Credit: u/rishabhsonker

Plain text files at the **user**, **project**, or **subfolder** level that Claude Code reads **automatically at the start of every conversation**.

They encode your conventions and constraints, and point to where the details live, so you do not have to repeat them each time.

The single most effective way to prevent errors before they happen.

Note: This does not guarantee Claude Code will always follow these instructions, but it significantly increases the chances.

What Belongs in CLAUDE.md

Every line is loaded in every session, so every line costs attention. **Keep a line only if Claude would work worse without it, and it will still be true next month.**

Belongs

- **Rules and prohibitions:** "never drop observations without logging," "all merges must report row counts," what Claude Code should **not** do
- **Conventions** that change the code it writes: language, libraries, paths, random seeds
- **Directory map:** where raw data, scripts, and outputs live
- **Commands** to run the pipeline and tests
- **Pointers** to the files below

Belongs elsewhere

- **Current status** ("debugging the merge"): stale within days, and then misleading
- **Variable definitions and codebooks:** long, needed only for some tasks
- **Research narrative and motivation:** written for human readers
- **What the code already says:** a second copy to keep in sync
- **Why you made past decisions:** needs dates, so it goes in a log

Keep *user-level* CLAUDE.md for things that apply across projects; use a *subfolder* CLAUDE.md only when a directory needs its own rules.

Which File Holds What

File	Who reads it	What goes in it
`CLAUDE.md`	Claude, *automatically*, every session	Rules, conventions, directory map, commands, pointers. Changes rarely
`README.md`	Humans new to the project: coauthors, replicators, you in a year	Question, data sources and access, how to reproduce, which output is which
`CODEBOOK.md`	Anyone touching a variable	Names, definitions, units, sources, construction rules
`STATUS.md`	You and Claude, at the start of the next session	Where you stopped, open problems, next tasks. Overwritten often
`docs/decisions.md`	Coauthors, referees, you at revision	Dated, append-only: what you decided, why, and the alternatives

CLAUDE.md is the only one Claude loads on its own

It reads the others when a task needs them or when you point it there. So CLAUDE.md stays short and mostly pointers, and each fact lives in **one** file, so the files cannot contradict each other.

CLAUDE.md Example

```
# CLAUDE.md - County Migration Project

## Layout and commands
- data/raw/: never edit. data/clean/: written by scripts
- src/: scripts, run in order. Run all: python src/run_all.py
- outputs/: tables and figures

## Rules
- Never drop observations without logging count and reason
- All merges must report pre- and post-merge row counts
- Always use robust standard errors
- Do not choose model specifications
- Do not interpret coefficients

## Read before changing things
- Variables and units: CODEBOOK.md
- Sample decisions: docs/decisions.md
- Where we left off: STATUS.md
```

Where CLAUDE.md Acts in the Loop

The same sequence as before, with one step highlighted.

Once, before you type anything

1 **Load** `CLAUDE.md`, your settings, and the tools it may use

The loop: repeats until the task is done

2 **Read** your prompt
3 **Gather** open files and search the project for what is relevant
4 **Reason** decide the next step
5 **Act** edit a file or run a command, asking your permission where required
6 **Observe** read the output or error that comes back, then return to step 3

When it stops

7 **Report** what it did and what it found

CLAUDE.md is loaded before your first message

In context from step 1, so you need not repeat it in every prompt. Claude can still misread it.

Exercise: Build a CLAUDE.md for Your Own Project

Use one of your own projects, in a fresh window so the project folder is the working directory.

1. **New VS Code window:** `Cmd/Ctrl+Shift+N`, or File → New Window
2. **Open your project root:** File → Open Folder
3. **Launch Claude Code** (`claude` in the terminal, or the Claude panel)
4. **Create** `CLAUDE.md` **at the project root**, with one of two routes:

Route A: Manually

- Explorer → New File, named exactly `CLAUDE.md`
- Write the layout, rules, and pointers, as in the example

Route B: With Claude

- Type `/init`: Claude reads your folder and drafts one
- Or ask for a draft with the layout, rules, and pointers
- **Read and edit it:** cut lines that fail the test

Check it worked: restart Claude Code, then type `/memory` to see the file listed, or ask "What do you know about this project?"

Session 1 Summary

1. Claude Code works with **files**, not interactive sessions
2. **You** set what it may do: permissions, and what stays out of the folder
3. What you get out depends on **your skill, the context, and the model**
4. The **context window** limits what it can hold
5. **CLAUDE.md** is your most powerful tool for preventing errors
6. **Cost management** is a core skill, not an afterthought

Next session: managing context, running sessions in parallel, plan mode, and verification

Keep your `my_workspace/session_1` **folder:** Session 2's first exercise builds on it.