

Session 2: Research Pipelines

From Data to Results

Brady Allardice

UPF & IBEI

Session 1: What Claude Code can do. How it works. CLAUDE.md and plan mode.

Session 2: How to use it on a real research task — and stay in control.

Today you will:

- ➊ Merge two datasets with explicit instructions
- ➋ Estimate a model and verify the output
- ➌ Generate robustness checks
- ➍ Produce a $\text{\LaTeX}$ results table

The methods are straightforward. The skill is the workflow.

Start-of-Session Ritual

Work on a copy, not inside the shared repo. Do this now:

Mac / Linux

```
# Already cloned? Update:
cd ~/claude-code-workshop && git pull

# First time:
git clone https://github.com/
  bradyallardice/claude-code-workshop.git

# Copy and open
cp -r session.2/ ~/my.workspace/session.2/
cd ~/my.workspace/session.2
code .
```

Windows (PowerShell)

```
# Already cloned? Update:
cd ~\claude-code-workshop; git pull

# First time:
git clone https://github.com/
  bradyallardice/claude-code-workshop.git

# Copy and open
Copy-Item -Recurse session.2 '
  ~\my-workspace\session.2
cd ~\my-workspace\session.2
code .
```

If `git pull` Fights Back

If you committed work inside the repo, `git pull` will complain. Run this recovery in the course repo, then continue with Steps 2–3:

```
git stash -u # save uncommitted edits
git branch backup-prework # save any local commits
git fetch origin
git reset --hard origin/student # match the server
git stash pop # skip if stash was empty
```

Your old commits live on `backup-prework`. If that name exists, add a suffix: `backup-prework-2`.

Strategies for Working Within the Context Window

Recall from Session 1: Claude Code can only “see” a limited amount of text at once. Errors compound when it loses track of earlier decisions.

Five strategies to work within this constraint:

- ① Keep tasks focused
- ② Maintain project documentation
- ③ Summarize work between conversations
- ④ Validate iteratively
- ⑤ Write token-aware prompts

What Happens When the Context Window Fills Up

As a conversation grows, Claude Code eventually runs out of room.

When this happens, it **auto-compacts**: it summarizes older parts of the conversation to free up space.

What auto-compacting does:


- Replaces earlier messages with a compressed summary
- Keeps recent messages and tool results intact
- Happens automatically. You will see a notice in the chat

The risk: Details from early in the conversation (variable definitions, decisions you discussed, constraints you mentioned) may be lost or simplified in the summary.

You can also compact manually:

Type `/compact` to force a summary and free up context space.

This is why CLAUDE.md matters: it survives compacting because



Strategy 1: Keep Tasks Focused

One topic, one conversation.

Avoid:

- “Clean the data, build the panel, run the regressions, and make tables”
- Long conversations that drift across topics
- Asking Claude Code to hold your entire project in its head

Instead:

- “Clean and reshape the raw census file”
- Start a new conversation for each distinct task
- Let each conversation do one thing well

Why

A focused task fits comfortably in the context window. A sprawling task forces Claude Code to forget earlier steps, and that is where silent errors enter.

Strategy 2: Maintain Project Documentation

Claude Code reads certain files **automatically** every time it starts.

Keep these up to date:

- `CLAUDE.md`: conventions, constraints, variable definitions
- `README.md`: project structure, what each directory contains
- Directory-level `CLAUDE.md` files for subfolders with specific rules

The Payoff

Good documentation means Claude Code starts every conversation already knowing your conventions, without you spending tokens re-explaining them.

Strategy 3: Summarize Work Between Conversations

When you finish a conversation, the context is gone. The next conversation starts fresh.

Before ending a conversation:

- 1 Ask Claude Code to summarize what it did and what remains
- 2 Save that summary (in a file, a commit message, or your notes)
- 3 Paste the summary into the *next* conversation as context

Example prompt:

```
Summarize what we did in this session: what files  
were changed, what decisions were made, and what  
still needs to be done.
```

Think of it like a handoff memo between research assistants.

The next “assistant” knows nothing unless you brief it.

Strategy 4: Validate Iteratively

The Compounding Error Problem

If step 1 is wrong and you do not catch it, steps 2–5 build on that mistake. By the time you check the final output, the error is buried.

Check each step before moving to the next:

- ① Clean the data → *verify row counts and distributions*
- ② Construct variables → *spot-check values against known cases*
- ③ Run the merge → *confirm pre- and post-merge row counts*

Slower per step, but **much faster overall** because you never have to unwind a chain of errors you didn't catch.

Strategy 5: Token-Aware Prompting

Every word in your prompt costs tokens. Make them count.

A good prompt tells Claude Code three things:

- ❶ **What files matter:** “Read data/raw/census.csv and scripts/clean.py”
- ❷ **What the goal is:** “Create a county-year panel with one row per county per decade”
- ❸ **What constraints apply:** “Keep all 3,143 counties; do not drop rows with missing population”

Vague:

“Clean up the data and make it ready for analysis”

Specific:

“Read census.csv, drop rows where year < 1950, and save as panel_clean.csv”

The Rule

Be specific about *inputs*, *outputs*, and *constraints*. Be brief about everything else. Context is finite; spend it on what matters.

The Core Workflow

Step	Who	If something is wrong
1. Describe	You	Rethink before sending
2. Review plan	You *	Reject, re-instruct
3. Generate & validate	Claude Code + You	Read the code, validate, re-instruct
4. Save or re-instruct	You	Keep the result (git in Session 3)

* = checkpoint where you can catch a problem before anything gets saved.

Validation is not a separate step — it happens **continuously** throughout steps 1–3.

The pattern

When something goes wrong, go back to **whichever step actually fixes the problem**: a bad plan means rewriting the instruction (step 1), not just re-running the same one. Diagnose first, then go back to the right step.

Step 1: Describe — Write It Down First

Before you type anything into Claude Code, write down your instruction.

If you cannot state the task precisely on paper, Claude Code cannot execute it precisely in code.

A good instruction specifies three things:

- ① **Inputs:** “Read `data/survey.csv` and `data/demographics.csv`”
- ② **Goal:** “Left-join demographics onto the survey by respondent ID”
- ③ **Constraints:** “Output must have exactly 2,044 rows. Do not drop respondents.”

Vague:

“Merge the datasets and run a regression”

Specific:

“Left-join demographics onto the survey by `respondent_id`. Use `control` as the reference. Run a weighted logistic regression.”

Be specific about inputs, outputs, and constraints. Be brief about everything else.

Step 2: Review the Plan

In plan mode, Claude Code proposes an approach before acting.

What to look for:

- Does it read the right files?
- Does the merge/analysis strategy match what you specified?
- Does it handle edge cases (missing data, extra rows, reference categories)?
- Does it make any *analytical decisions* you did not ask for?

The key question

Is there anything in this plan that I did not explicitly ask for?
If yes: reject or modify before it runs.

You do not always need plan mode. For simple tasks — summary statistics, reformatting, boilerplate — steps 1–3 collapse into one. But when uncertain, *use it*.

Step 3: Generate and Validate

Claude Code writes and runs the code. Your job is not done.

Read what Claude Code produced:

- Does the script do what the plan said it would?
- Are variable names and file paths correct?
- Did it make any choices the plan did not mention?

Common surprises at this stage:

- Silently dropping rows with missing values
- Choosing a default reference category you did not specify
- Using unweighted estimation when you asked for weighted
- Adding “helpful” data cleaning steps you did not request

The habit

Read the code, not just the output. Output can look correct even when the code is wrong.

Validation Strategy 1: Read the Code

Claude Code is not deterministic — the same prompt can produce different code each time. But **code itself is deterministic**. If you read the code and understand what it does, you can be confident in the results, the same way you would be with code you wrote yourself.

This is the most fundamental verification:

- Read the script Claude Code produced, line by line
- Confirm it does what you asked — not more, not less
- If you understand the code and it is correct, the output is correct

The principle

You are not trusting Claude Code. You are trusting *the code* — the same way you trust code you wrote yourself. The difference is who typed it, not whether it works.

Validation Strategy 2: Check Known Properties

Before you look at results, ask: *what do I already know must be true?*

Examples:

- A left join should not change the number of rows in the left table
- Adding a control variable should not change the sample size (unless it has missing values)
- A share variable must be between 0 and 1
- Treatment groups should sum to the total sample

Ask Claude Code to assert these explicitly:

- “Assert the merged output has exactly 2,044 rows”
- “Print the number of observations used in each regression”
- “Verify no duplicate respondent IDs”

If the assertion fails, you catch the error immediately — not three steps later.

In VS Code: When working directly in the editor, you see the `git diff` before you accept the change. This shows exactly what Claude Code changed.

Validation Strategy 3: Write Down Expectations First

The Pre-Registration Analogy

Just as you pre-register hypotheses before seeing results, write down what a correct transformation should produce *before* you run it.

Before a merge or transformation, state:

- Expected output dimensions (rows $\times$ columns)
- Which columns should appear and their types
- How missing values should be handled
- A specific row you can trace through by hand

Then compare the actual result to your expectations.

Discrepancies are either a bug or a misunderstanding — both worth catching early.

Validation Strategy 4: Unit Tests Against Known Values

Pick a handful of observations and compute the correct answer by hand.
Then verify the code reproduces them.

Example prompt:

Write a test that verifies:

- Respondent 9 (interviewed 2015-10-07) has
treatment = "history" and fx_status = "none"
- The control group has exactly 480 respondents
- The weighted proportion supporting intervention
is approximately 0.486

Why this works:

- Forces you to engage with the data at the observation level
- Catches silent recoding, wrong joins, and off-by-one errors
- Tests survive across refactors — rerun them after every change

Validation Strategy 5: Replicate Before You Extend

Building on someone else's work?

Have Claude Code reproduce their results first.

- ① Give Claude Code the original paper's code or description
- ② Ask it to replicate a key table or figure
- ③ Compare output to the published result
- ④ Only after replication succeeds, begin your extension

Why

If you cannot reproduce the baseline, anything you build on top of it is uninterpretable. Replication is not a detour; it is the foundation.

Validation Strategy 6: Use a Second Agent as a Critic

Use a *second* Claude Code conversation as a reviewer.

How:

- 1 Agent 1 writes the code
- 2 Open a new conversation (Agent 2) and ask it to review:
“Read the analysis script and identify any bugs, silent data loss, or decisions that could change the results.”
- 3 Agent 2 has fresh context and no attachment to the code

Why this is powerful:

- The writing agent has blind spots — it wrote the code and “believes” it is correct
- A fresh agent reads the code like a skeptical reviewer
- Catches issues you would miss because you watched it being built

Treat your agents like co-authors who check each other's work, not a single assistant you trust unconditionally.

Step 4: Save or Re-instruct

If the output passes verification: keep it. Claude Code saves its changes to your files after each step, so the results are already on disk.

If something is wrong: identify the specific error and re-instruct. Do not ask Claude Code to “fix it” — tell it *what* to fix and *how*.

In Session 3 we will learn **git** — a tool that lets you snapshot your project before each change and revert if anything goes wrong. For now, know that the save step exists and that there is a better way to manage it coming soon.

The full cycle

Describe → Review plan → Generate & validate → Save or re-instruct

Repeat for each task in your analysis. One cycle per task, not one cycle for the whole project.

What to Delegate, What to Verify, What to Keep

Delegate freely

Boilerplate & docs
File management[†]
Reformatting code

Delegate + verify

Data loading & cleaning
Data merges
Variable construction
Summary statistics
Robustness variants
Visualization
L^AT_EX tables

Do not delegate

Choosing specifications
Interpreting results
Identification strategy
Sample selection
Causal claims

[†] Be careful with file management — Claude Code can delete files. Always review before approving.

The test: *If a referee questioned this choice, who would need to answer — you or the RA?*

Today's exercise focuses on the middle column: tasks where Claude Code is productive but where you must verify every output.

Today's Data: The Swiss Franc Shock

Background: In January 2015, the Swiss National Bank removed its currency peg. The Swiss franc surged against the Polish zloty. Polish households with CHF-denominated mortgages saw their payments spike overnight.

The survey: 2,044 Polish adults, interviewed October 2015. Respondents were randomly assigned to receive different information about the crisis before being asked whether the government should intervene to help affected borrowers.

Your files:

- `swiss_franc_survey.csv` — 2,044 respondents, 11 variables
- `respondent_demographics.csv` — income and ideology data (more rows than the survey)
- `codebook.md` — variable descriptions

Source: Ahlquist, Copelovitch, and Walter, "The Political Consequences of External Economic Shocks: Evidence from Poland," *American Journal of Political Science* (2020).

Does providing information about the Swiss franc crisis increase support for government intervention to help affected borrowers?

Treatment (randomized):

- Control (no info)
- Factual information
- Historical context
- Hungary's policy response

Outcome:

- Supports government intervention (binary)

Moderator:

- FX loan exposure status

The methods are straightforward. The skill is staying in control of the workflow.

Exercise Part 1: Merge and Explore (30 min)

Use plan mode for every step. Review every plan before approving.

- ❶ **Merge** the demographics data into the survey.
 - The demographics file has more rows than the survey
 - **You decide the join type.** Do not let Claude Code choose.
- ❷ **Explore** the merged data: summary statistics, treatment group sizes, outcome distribution.
 - **Before you look at results: write down what you expect to see.**
- ❸ **Verify** the merge.
 - What should you check? Write down your verification checks *before* running them.
 - Then run them.

Part 1 Debrief: What Did You Verify?

What did you check?

What did you check?

You should have verified:

- Row count: exactly 2,044 rows (no rows gained or lost)
- Join type: only a left join gives the correct result
- No duplicate respondent IDs in the output
- Treatment groups: 480 / 549 / 520 / 495
- Proportion supporting intervention: ~ 0.479 (unweighted)
- FX exposure: 1,918 / 69 / 55

Did your numbers match? Did Claude Code choose the right join type?

Exercise Part 2: Specify and Estimate (30 min)

- 4 **Base model:** Logistic regression of `supports_intervention` on `treatment`.
 - Use survey weights
 - What should the reference category be? Tell Claude Code explicitly.
- 5 **Add FX exposure.** Did the sample size change? Why or why not?
- 6 **Add demographic controls:** age, female, education, urban/rural, income, left-right.
 - Before running: predict how many observations you will have.
 - Then check. Were you right?
- 7 **Add an interaction** between `treatment` and `FX exposure`.

Part 2 Debrief: What Happened to Your Sample?

How many observations did you have in each model?

Part 2 Debrief: What Happened to Your Sample?

How many observations did you have in each model?

The Missing Data Trap

Adding control variables can silently change your sample.

Specification	N	What happened
Treatment only	2,036	Full sample (8 missing outcome)
+ FX exposure	2,035	2 missing FX status
+ age, female, education, urban/rural	2,035	No additional missingness
+ income quintile	~1,473	565 missing income
+ left-right scale	~1,499	542 missing left-right
+ both	~1,151	44% of sample dropped

Did Claude Code mention this when it added the controls?

Exercise Part 3: Robustness and Output (30 min)

- 8 **Generate robustness checks** from your base model:
 - With and without survey weights
 - Different control sets (track how N changes each time!)
 - Collapse three treatment arms into “any information” vs. control
 - Restrict to complete cases so N is constant across specs
- 9 **Format as $\text{\LaTeX}$.** Ask Claude Code to produce a publication-style regression table with all specifications side by side. Compile to PDF.
- 10 **Commit.** Save all scripts, output, and the $\text{\LaTeX}$ table.

You specify the core analysis.
Claude Code generates the variations.
You verify each one.

Claude Code never chooses what robustness checks to run.

For each specification, check:

- Does it use the right sample? (What is N?)
- Does the code match what you asked for?
- Are the coefficients in a plausible range?
- Do the signs on control variables make sense?

L^AT_EX Tables: Bringing Results Into Your Paper

Until now, your results are printed to the console or saved as CSVs. In your actual workflow, they need to end up in a paper.

Why this matters:

- Claude Code can write L^AT_EX tables directly from regression output
- The table lives in your project directory, under version control
- When you update the analysis, Claude Code updates the table
- No manual reformatting, no copy-paste errors

Today: generate 5 tables instead of one table. In Session 3, we bring the full paper into VS Code.

The bigger idea

The more of your research that lives inside the project directory, the more Claude Code can operate on. Code, data, output, tables — and eventually the paper itself.

What to Watch For During the Exercise

This exercise is designed to surface moments where Claude Code makes choices for you.

- **The merge:** Did Claude Code pick the right join type? Did it mention the extra rows?
- **Reference categories:** Did it set treatment to “control” or pick alphabetically?
- **Missing data:** When you added controls, how many observations did you lose? Did Claude Code mention this?
- **Survey weights:** Did it use them unless you told it to?
- **Income = 0:** “No income” is coded as 0. How did Claude Code treat it?

These are the moments where your judgment matters.

The workflow is: notice the decision, direct Claude Code explicitly, then verify the output.

Discussion questions:

- 1 Did your merge lose any observations? How did you handle weekends?
- 2 How many observations did you lose when you added controls?
Did Claude Code tell you?
- 3 How many of your robustness specifications were correct on the first try?
- 4 What did you catch in the plan that you would have missed in the output?

Key takeaway

The assertions and the plan review are not extra work — they are the workflow. The 30 seconds it takes to read a plan or check a sample size can save hours of debugging downstream.

Session 3: Code Auditing, Git, and $\text{\LaTeX}$ in VS Code

- Claude Code reads and critiques *your* code instead of writing it
- Git as your safety net: commit $\rightarrow$ instruct $\rightarrow$ diff $\rightarrow$ revert
- Bringing your paper into the project: $\text{\LaTeX}$ in VS Code

Before next session:

- Try the core workflow on something from your own research
- Note where Claude Code made a decision you did not ask for
- Install the $\text{\LaTeX}$ Workshop extension in VS Code (instructions in setup guide)

