

Session 4: Advanced Claude Code

Skills, MCPs, and Subagents

Brady Allardice

UPF & IBEI

Where We Are

Session 1: What Claude Code can do. How it works. CLAUDE.md and plan mode.

Session 2: The core workflow. Data cleaning, estimation, robustness, $\LaTeX$ tables.

Session 3: Git and $\LaTeX$ in VS Code. Your paper joins the project.

Session 4: The layer above the core workflow — infrastructure.

Today you will:

- ➊ Migrate to a cleaner repo workflow (a structural fix, not an agent fix)
- ➋ Learn to route a wish — is it a **skill**, **project context**, or **MCP**?
- ➌ Connect two MCPs (Zotero, GitHub) in practice
- ➍ See when to reach for a **subagent** — and when not to

Today is about knowing which tool does real work and which is just novelty.

The Dangers of Emerging Technologies

Claude Code's prying AIs read off-limits secret files

Developers remain unsure how to prevent access to sensitive data

Wed 25 Jan 2025 - 20:42 UTC

Thomas Claburn

Developers told me that Claude Code was following simple instructions? Anthropic's Claude Code continues to read passwords and API keys, even when blocked.

passwords, tokens, API keys, and other sensitive files. And if they do so, they're supposed to be blocked in a publicly accessible .git repository.

With a .gitignore file that tells the system to ignore copying a local repo to a remote server, the .gitignore file.

When you ask Claude to read the contents of a file in your project root. This works like a charm for files matching patterns listed there."

But in a recent post, Claude can read the contents of a .gitignore file that ought to prevent

- Thought for 0s >
- Read .claudeignore
- Read .gitignore

- Thought for 3s >

- You've found the problem! .claudeignore correctly lists .env as off-limits with exceptions for .env.test and .env.example:

```
.env      - Should be protected
.env.*    - All variants should be
!.env.example - Exception (but we delete it)
!.env.test  - Exception (allowed for testing)
```

The issue: I was able to edit .env despite it being in .claudeignore. This is a security gap—the .claudeignore file should have blocked that write. The .claudeignore file has the right config, but the Claude Code system didn't enforce it.

[BUG] .claudeignore does not prevent Claude from reading ignored files #36163

worth knowing - you can add a .claudeignore file to your project root to tell Claude Code to skip those files entirely. so .env, *.pem, secrets/ etc can all be excluded. not a perfect solution since it's opt-in per project, but it's the built-in way to handle this without needing extra tooling.

thesaadmirza OP · 15d ago

Yeah .claudeignore helps with the discovery side. Worth using.

Claude Ignore

Similar to .gitignore, there is a .claudeignore file. Entries in the .claudeignore tell Claude that the files and folders listed are not relevant.

Folders like node_modules, vendor, logs, caches etc would be good one to add to .claudeignore.

.claudeignore is not a replacement for settings.json. Claude can still read items listed in .claudeignore. You can still ask Claude to read them to gain context.

.claudeignore saves you tokens by not having Claude read the files upfront.

Vista creada con IA

The .claudeignore file is an official, recommended mechanism for Claude Code (Anthropic's agentic CLI tool) to define which files and directories the AI agent should skip reading or indexing. [socket.dev v1](#)

It serves as a guardrail to maintain privacy and reduce unnecessary context consumption, similar to how .gitignore operates. [socket.dev v1](#)

`.claudeignore`: Works, Sometimes

The `.ignore` convention is everywhere in software:

- `.gitignore` (git), `.dockerignore` (Docker)
- `.npmignore` (npm), `.eslintignore`, `.prettierignore`
- `.gcloudignore`, `.terraformignore`, and many more

So `.claudeignore` is a commonly taught, perfectly reasonable extrapolation. You'll see it recommended across blogs, Reddit, and LLM-generated advice.

What the testing actually shows:

- If you tell Claude *explicitly* to respect it, it usually does.
- If you don't — even after it reads `.claudeignore` and sees the pattern — it will read the file anyway.
- There's no tool-level enforcement. It's a polite request that only lands when the model is listening.

The Official Approach — With Real Gaps

The docs point you at `permissions.deny` in `.claude/settings.json`:

```
{
  "permissions": {
    "deny": [
      "Read(session_4/demo_output/tab_policy.*)",
      "Grep(session_4/demo_output/tab_policy.*)",
      "Glob(session_4/demo_output/tab_policy.*)",
    ]
  }
}
```

The catch: it blocks Read/Grep/Glob, not Bash. A `cat` in a shell command reads the file anyway.

Safe assumption: if it's in the project, Claude can read it. Be deliberate. Stronger mechanism coming with **hooks**.

Since Session 3:

- Did the git cycle (commit → instruct → diff → decide) stick?
- Did anyone bring their paper into VS Code? How did it feel?
- Any tool fight back at you — compilation errors, merge conflicts, something else?

Hold onto the merge-conflict stories. We start there.


First: What Is GitHub?

GitHub is where code lives on the internet.

Every project is a *repository* — a folder with all its files plus the full history of every change.

Two copies of everything:

- **Remote** — the shared version on `github.com`
- **Local** — the version on your machine



The short version

GitHub hosts the *shared* copy. You work on your *local* copy. Git keeps them in sync.

The Vocabulary

Five words that will keep coming back today:

- **clone** — make your first local copy of a remote repo
- **pull** — fetch the latest changes *from* the remote into your local copy
- **push** — send your local changes *to* the remote
- **branch** — a parallel line of work inside the repo; changes on a branch don't affect anyone else until it's merged
- **pull request (PR)** — a proposal to merge one branch into another, with a page on GitHub for review and discussion

The flow

clone once. pull before you start. Make a **branch**, commit, push, open a **PR**. Someone reviews and merges.

The Course Repo

Our materials live in one GitHub repo. Everything I push goes there; everything you pull comes from there.

Let's look at the page together. Things to notice:

- The **file tree** — `session_1/`, `session_2/`, ..., `session_4/`
- The **README** — what the repo is and how to use it
- The **commit history** — every change I've made, with a message
- The **branches** — `main` is the shared one; others are work-in-progress

Why this matters for what's next

The issue we hit last session lives exactly here: you and I were both editing files in the *same* repo. To explain what went wrong, we need the picture of where things live.

The Problem We Had

Every session so far, someone has hit a merge conflict.

What we did about it:

- Deleted the local repo and recloned
- Asked Claude Code to resolve the conflict markers
- Copied work to the Desktop, wiped the directory, copied back

That worked. Everyone got unstuck. But it was a symptom fix.

The Root Cause

You were editing files *inside* the shared course repository. Every time I pushed new material, git had to reconcile your changes with mine. Conflicts were guaranteed by the structure, not by accident.

The Structural Fix

New layout on my side: each session is a self-contained folder.

```
AIAgentsCourse/  
  session_1/  
  session_2/  
  session_3/  
  session_4/ <- today's materials live here  
  ...
```

New workflow on your side: copy the session folder out of the repo into your own workspace. Work there. The shared repo stays read-only for you.

What this changes

I can push new material anytime without breaking anything in your workspace. You can edit, break, delete, and redo without worrying about the shared state. No more conflicts.

The Three-Step Ritual

Do this at the start of every session from now on:

```
# 1. Pull the latest course repo
cd ~/AIAgentsCourse && git pull

# 2. Copy today's session folder into your workspace
cp -r session_4/ ~/my_workspace/session_4/

# 3. Open your copy in VS Code and work there
cd ~/my_workspace/session_4 && code .
```

We do it together now, live. Five minutes. Everyone ends up in the same clean state before we move on.

If you had in-progress work inside the course repo, copy it into `~/my_workspace/` too. I will help you sort out what to keep.

Claude Skills

A **skill** is a reusable instruction bundle
Claude loads on demand when a task matches.

It codifies a *convention* —
how *you* do a thing correctly —
so you don't re-explain it every session.

What a Skill Actually Does

Skills from my setup, and when each one fires:

- `latex-regression-table` — I ask for a results table. It writes a canonical CSV, renders a booktabs `.tex`, and compiles a PNG. One house style across the whole paper.
- `abstract-structure-audit` — I paste a draft abstract. It checks whether the four standard moves (puzzle, approach, findings, contribution) are present, vague, or missing.
- `paragraph-structure-audit` — I share a paragraph. It flags multiple claims per paragraph, smuggled assumptions, and weak transitions.
- `robustness-checks` — I run a regression. It proposes a structured battery (alternative SEs, subsamples, IV, event study) and executes it.

Each one is a recurring workflow done the same way every time.

How a Skill Works

A skill is a folder with a `SKILL.md` file in it.

What you set up:

```
~/.claude/skills/latex-tables/  
SKILL.md <- required  
template.tex <- optional  
examples/ <- optional
```

`SKILL.md` starts with frontmatter:

```
---  
name: latex-tables  
description: Use when generating a  
  regression or summary-stats table.  
---  
  
# Instructions Claude follows  
...
```

How Claude uses it:

- 1 At session start, Claude sees every skill's name + description — nothing more
- 2 **Two ways to invoke:** auto-call on a description match, or you call it yourself (`/latex-tables`)
- 3 The full `SKILL.md` body loads into context on demand
- 4 Claude follows those instructions

The key idea

The description is the auto-call trigger. Write it so Claude can tell *when* this skill applies.

Before You Build a Skill, Ask Where It Belongs

Three things look alike but are different tools:

- ① **Skills** — reusable instruction bundles, loaded on demand
- ② **Project context** — a CLAUDE.md inside your paper's directory
- ③ **MCPs** — connections to external systems

Most wishes of the form “I wish Claude knew X” route to one of these three.

Knowing which one saves you from building a skill that should have been something else. It also tells you when a skill is *too narrow* a tool for the job.

Skills → what's stable *about you*

Projects → what's stable *within a paper*

MCPs → what's *live*

If X changes when the project changes, it's not a skill.

If X is alive in an external system, it's not a skill either.

The Rule, Applied

What you want Claude to know	Where it belongs
Your $\text{\LaTeX}$ table conventions	Skill (stable about you)
The variables & specs of your current paper	Project context (within a paper)
Your Zotero library	MCP (live external state)
Your writing voice	Skill (stable about you)
What your current paper argues	Project context (within a paper)

Why this matters

Skills are a narrower tool than most people assume. That's a feature — narrower tools are easier to use well.

Within Skills: Two Different Jobs

Verification skills

Fire *during or after* work to check it.

Examples: LLM-mistake checker, sanity-check validator, replication auditor.

Earn keep by: catching real problems, even if used infrequently.

Production skills

Fire *during the making* of something to shape it.

Examples: $\text{\LaTeX}$ table skill, writing-style skill.

Earn keep by: frequency — every table, every document.

Why the distinction matters

“Is this worth it?” has a different answer for each. Verification skills need to catch high-cost errors. Production skills need to fire often.

Skills are for **conventions**
you hold about how to do
a thing correctly.

Conventions about table formatting. About sanity-checking data.
About responding to reviewers. About engaging with the literature.

If what you want to codify isn't a convention, it's probably not a skill.

Two skills I reach for regularly:

1. **L^AT_EX tables.** Encodes my house style for regression tables: significance notation, SE labels, table-notes format, journal overrides. Fires every time I generate a table. Saves ~15 min per table and keeps formatting consistent across a paper. On version three.
2. **Paragraph-structure audit.** One claim per paragraph, mis-cited claims, smuggled assumptions, weak transitions. Fires when I'm editing my own drafts — catches the kind of sloppiness I stop seeing after four re-reads.

What they have in common

Each codifies a *convention* I already hold. Each fires at a predictable moment — generating a table, reviewing a draft. Each has been revised through actual use.

Demo

The $\text{\LaTeX}$ tables skill, applied live.

Finding Good Use Cases Is Hard

Two skills I built but barely use:

- `auto-commit` — commits after every file edit
- `spec-validator` — pre-estimation data and spec checks

What they have in common: each sounded useful in the abstract. Each was built on an idea of how research *should* go. Each got replaced by a one-off prompt the first time I actually needed the job done.

The specificity can also work against you — a tightly scoped skill only fires when the exact shape of the task matches, and it's easy to build one whose niche never actually shows up.

Strong recommendation

Do not go build five skills this weekend. You will not use most of them.

Instead: **build one**. Try it on your current work. Tweak it where it fires wrong. Let it earn its place in your workflow. *Then* add another.

How to Build One That Works

What separates a skill that fires right from one that sits unused:

- **Write the description for the trigger, not the reader.** Claude decides whether to auto-call on the description alone. Name the task, the keywords, and what the skill does *not* handle.
- **Keep it narrow.** One genre of table, one kind of audit, one workflow. Broad skills fire at the wrong moment, or never.
- **Encode your convention, not generic knowledge.** The SKILL.md should capture the judgment calls *you* make. Claude already knows boilerplate.
- **Ship a reference example.** An examples/ folder with a real working artifact gives Claude something concrete to pattern-match on.
- **Revise after the first real use.** Version one fires wrong somewhere. Fix the trigger or the instructions, retry.

Let's Build One: tex-reference-audit

The need: before submission, I want to know every `\cite`, `\ref`, and `\label` in the manuscript actually resolves.

What it checks:

- Every `\cite{key}` resolves to a `.bib` entry
- Every `.bib` entry has required fields (author, year, title, journal, pages...)
- Every `\ref{label}` points to an existing `\label`
- Every `\label` is actually referenced (no orphans)
- Figures use `fig:`, tables use `tab:`

The description — copy, paste, build live:

```
---
name: tex-reference-audit
description: Use when reviewing
  or finalizing a LaTeX manuscript.
  Audits \cite against the .bib
  file, \ref/\label pairs for
  figures and tables, flags
  missing fields, duplicates, and
  orphans. Trigger: "check refs,"
  "audit citations," "review
  LaTeX."
---
```

No script — the body is just instructions Claude follows to parse, check, and report.

Exercise: Build a Skill (25 min)

Working solo. Each of you walks out with one skill installed and tested.

- ➊ **Pick a convention (3 min).** Something you actually did in the last two weeks — a format you enforce, an audit you redo by hand, context you retype. Not a hypothetical.
- ➋ **Route it (2 min).** Confirm with your neighbor: is this *really* a skill? Not project context, not an MCP. If a skill, is it **verification** or **production**?
- ➌ **Write SKILL.md (10 min).** A clear name, a description that tells Claude *when* to fire, and the instructions for when it does.
- ➍ **Install (1 min).** Save to `~/.claude/skills/<your-skill>/SKILL.md`.
- ➎ **Try it (9 min).** Open a *fresh* Claude Code session. Ask for a task that should trigger it. If it doesn't fire, the description is wrong — fix and retry.

A working skill, installed and fired at least once.

What you leave with:

- A folder at `~/.claude/skills/<your-skill>/` with a real `SKILL.md`
- One transcript where Claude triggered the skill on a real task
- One concrete thing you'd change now that you've seen it run

Two or three people share: *What is it? Did it trigger on the first try? What surprised you?*

The goal

You leave having built one thing, seen it fire, and learned where the description is too vague or too narrow. That is the empirical loop. Everything else is speculation about skills you haven't tried.

You now have everything you need to make skills work for you:

- ① A **routing rule** to spot a real candidate — stable-about-you, not within-a-paper, not live
- ② A **typology** — verification vs. production — to predict when and how often it'll fire
- ③ A **heuristic** — skills codify the conventions you already hold
- ④ An **empirical loop** — build, fire, refine, which you just ran end-to-end
- ⑤ **One skill already installed** on your machine

The durable takeaway

Every convention you hold about how to do a thing correctly is a candidate skill waiting to be written down. The ones that earn their place will compound across every project you work on — and you now know how to tell the good candidates from the noise.

Connecting Claude Code to systems outside your project.

From the routing rule: *MCPs are for what's live.*
Here are two live things worth connecting.

Forty-five minutes. Two MCPs. One pair exercise.

What an MCP Is

MCP = Model Context Protocol. An open standard for plugging external tools into Claude Code.

- Each MCP is a separate server.
- Claude Code speaks to it over a standard protocol.
- The server exposes *tools* (like `zotero_search`) that Claude Code can call.

Without MCPs: Claude Code reads and writes files in your project. That's it.

With MCPs: it can query your Zotero library, open a GitHub PR, read a database, post to Slack — anything someone has built a server for.

The framing

MCPs turn Claude Code from a tool that operates on your *files* into a tool that operates on your *workflow*.

Where Skills and MCPs Compose

The paragraph-structure skill is the first one you've seen that invokes an MCP.

Alone:

- *Skill only*: "This claim needs a citation. Figure one out."
- *MCP only*: "Here's your Zotero library." But nothing knows when to search.

Composed:

- The skill knows *when* a claim is mis-cited or uncited
- The MCP returns actual papers from your library that would support it
- The MCP constrains the skill to real papers — no hallucinated citations

The pattern

A skill knows *when* to act; an MCP provides *live data*. Together, they do what neither can alone.

The Two-Reason Rule

Once you've routed something to an MCP, ask: which of two reasons?

- ❶ **Extend capability.** Let Claude Code do something it couldn't do at all.
Example: open a GitHub pull request. Query your Zotero library. Run a SQL query against your database.
- ❷ **Reduce friction at a boundary.** Let it do something you *could* already do, but where context-switching between tools kills the flow.
Example: inserting citations while drafting. You could look them up in Zotero and paste. But stopping to do that breaks writing flow.

Today I'm showing two specific MCPs I think are worth adopting. There are dozens more. The *framework* is what transfers; the two cases are how you learn it.

Callback to the meta-lesson: MCPs are easy to accumulate. Ask what each one earns.

The Friction Case: Zotero

Drafting with citations normally looks like:

- ① Write a sentence that needs a citation
- ② Stop. Open Zotero. Search. Find the paper. Copy the cite key.
- ③ Come back to the editor. Paste. Context-switch back into drafting.
- ④ Repeat ten times per paragraph.

With the Zotero MCP: “Find me the Smith 2019 paper on voter ID, add it to refs.bib, and cite it here.” Claude Code does the search, the BibTeX export, the `\cite{}` insertion — without you leaving the `.tex` file.

Why it earns its keep

Reduce friction at a boundary. You could already do all of this manually. The MCP just removes the tool-switch.

What the Zotero MCP Can Do

The Zotero MCP exposes around two dozen tools. They fall into two buckets:

Read from your library

- **Search:** keyword, semantic (natural language), by tag, advanced filters
- **Browse:** collections, recent items, saved searches, RSS feeds
- **Pull metadata:** title, authors, abstract, BibTeX fields
- **Pull full text** of attached PDFs
- **Read** your existing notes and annotations

Write back to your library

- **Add papers** by DOI, URL, or local PDF (auto-fetches metadata + open-access PDFs)
- **Create notes** attached to items
- **Create and organize collections**
- **Update item metadata** — tags, titles, abstracts

What this unlocks

Claude can read the *actual full text* of papers in your library — not just the title — so it can reason about what's *in* the paper, not just whether the citation exists.

A single prompt:

```
Find all papers in my aina_articles collection on  
local economic voting, and for each give me the  
abstract and any notes I have on it.
```

What you'll see in the tool-call log:

- ❶ `zotero_get_collection_items("aina_articles")` — lists the collection
- ❷ `zotero_semantic_search("local economic voting")` — filters to the topic
- ❸ `zotero_get_item_metadata(...)` — pulls title, authors, abstract
- ❹ `zotero_get_notes(...)` — pulls my existing notes on each item

Watch the tool calls, not just the output. Seeing which tools fire tells you what the MCP actually does.

Exercise: Mini Literature Review (15 min)

You installed the Zotero MCP as Session 4 prework and pointed it at the shared group library `aina-articles` — recent papers on economic voting.

Your task. Draft a **five-sentence literature review** on *one* of:

- How economic conditions shape vote choice
- Whether voters punish incumbents for bad economies
- How economic voting differs across contexts (regimes, institutions, crises)

Rules:

- Every sentence cites at least one paper from `aina-articles`
- Only use cite keys that *actually exist* in the library — no hallucinations
- Claude does the searching, full-text reading, and drafting

Prompt to start from:

```
Using only papers in my Zotero 'aina-articles' collection, find
5 relevant to [topic]. Pull the key claim from each full text,
then draft a 5-sentence lit review with \cite{} for each.
```

What to Watch For

While the model works, watch the tool calls:

- Did it run `zotero_semantic_search`, or skip search and guess?
- Did it actually pull `zotero_get_item_fulltext`, or cite from titles alone?
- Do the claims in your sentences match what the papers actually say?

After drafting, spot-check one citation. Open the paper in Zotero. Does the sentence you wrote accurately represent what the paper argues?

The real lesson

An MCP gives Claude *access*. It does not give Claude *judgment*. The failure mode to expect: a plausible sentence citing a paper that says something adjacent — but not quite that. **You still read.**

The Capability Case: GitHub

GitHub MCP lets Claude Code do things you couldn't do by just editing files:


- Open a pull request from a branch you've been working on
- Leave review comments on a coauthor's PR
- Create issues to track tasks you don't want to lose
- Push files to a remote repository

Why now? This morning's Git restructuring put you in a clean state. You are working in your own workspace. You understand the three states: local copy, pushed branch, shared main. The GitHub MCP is what closes the collaboration loop.

Why it earns its keep

Extend capability. "Open a PR" isn't something Claude Code can do without this. It's a genuine new capability, not a friction fix.

The Mental Model



Branches + PRs are the collaboration mechanism. The GitHub MCP is a *helper* on top — not a substitute.

The discipline

Use `git` directly. Use the MCP when you're stuck or when the MCP genuinely saves steps. When things break, you need to understand what's actually happening — and the MCP abstracts that away.

Exercise: The Collaborative Story

A shared class repo. You will write one chapter of a weird, episodic story. Your partner reviews and merges. At the end, we assemble all chapters and read the result out loud.

Three premises — we vote on one:

- ❶ **The Conference From Hell** — every session, break, and hallway moment at an academic conference goes sideways
- ❷ **The Dissertation That Wrote Itself** — a PhD student's dissertation has gained sentience; each chapter is one night of it doing something weirder
- ❸ **The Department Mouse** — a very literate mouse is stealing specific things from the department; it appears to have a plan

The writing is short on purpose. The point is the Git workflow, not the prose.

Your Role, Your Chapter

Pairs are pre-assigned. Find yours in `PAIRS.md` at the root of the class repo.

Each pair writes one chapter. Two roles:

- **Drafter** — writes the chapter, commits, pushes, opens the PR
- **Reviewer** — reads the partner's PR, approves, merges

Your chapter number is next to your name. The prompt for that chapter is in the winning premise file under `premises/`.

Ground rules

- **Length doesn't matter.** A paragraph is fine.
- **Don't try to match other chapters.** You can't — everyone is writing simultaneously.
- **Don't polish.** Five minutes. Write, push, move on.
- **Weird is the point.**

Drafter

```
git checkout main && git pull
git checkout -b chapter-N
# edit chapters/chapter_N.md
git add chapters/chapter_N.md
git commit -m "Draft chapter N"
git push -u origin chapter-N
```

Then open the PR:

- Click the banner on GitHub, base main, assign partner
- Or ask Claude Code: *"Open a PR from chapter-N to main"*

Time budget (~15 min): vote 1 | draft 5 | PR + review + merge 4–6 | assemble 2

Reviewer

- ① Wait for your partner's PR link
- ② Open it, click **Files changed**
- ③ Skim — if it reads like a chapter, done
- ④ **Review changes** → **Approve**
- ⑤ **Merge pull request** → **Confirm**

You are not grading. Approve and merge.

What to Expect

Some pairs will finish. Some will not. Both are fine.

The goal is that *everyone sees the loop end-to-end*: branch, push, PR, review, merge. Even if you need help at every step, you will recognize the shape next time.

What will go wrong (for some of you):

- **Push rejected** because `main` moved. Run `git pull --rebase origin main` and push again.
- **PR opened against the wrong base.** Edit the base branch in the GitHub UI, or close and reopen.
- **GitHub MCP asks for a token** you haven't set up. Use the GitHub UI — both paths work.
- **You get stuck for more than 90 seconds.** Flag me. Don't silently spin.

After all PRs are merged, I'll run the assembly script and read us the story.

GitHub MCP is optional infrastructure. The underlying git workflow is not.

If you find yourself reaching for the MCP to fix a conflict — sometimes that's the right move. But sometimes (like this morning's session-folder fix) it's papering over a structural problem.

Before adopting an MCP, ask

- 1 Does this genuinely extend what I can do, or reduce real friction?
- 2 Will I invoke it more than three times a month?
- 3 Am I using it because it helps, or because it's new?

Most MCPs fail question 2. The two today pass all three — *for me*. Your answers may differ.

Know they exist.

Know when to reach for them.

Don't over-invest.

Ten minutes. No hands-on.

What a Subagent Is

Claude Code can spawn another Claude Code instance inside your current session — a subagent.

- It has its own context window
- It gets a scoped task and a scoped tool set
- It returns a summary (not its raw output) to your main conversation

Two values:

- ① **Parallelism** — run three searches at once instead of serially
- ② **Context hygiene** — a big messy search doesn't pollute your main conversation

Think of it as

Delegating a bounded sub-task to a fresh colleague who doesn't see your notes — and who reports back only the conclusion.

When They Shine — and When They Don't

Subagents help when subtasks are:

- **Parallelizable** — independent work that doesn't need to be sequenced
- **Bounded** — clear inputs, clear outputs, no back-and-forth
- **Independently verifiable** — you can check the result without needing the agent's full trace

Research workflows that fit:

- Multi-database literature searches (one agent per database)
- Coding many documents against a fixed rubric
- Parallel robustness checks on independent specifications

Research workflows that don't fit:

- Anything sequential and judgment-heavy — most of your drafting, most of your analysis decisions
- Tasks that depend on your ongoing conversation context (subagents start cold)

Most research tasks do not have the shape that makes subagents worth it.

- Most of what you do is sequential
- Most of your decisions depend on context the subagent won't see
- Most of your verification needs the full trace, not a summary

If your workflow has parallel, independent, verifiable subtasks — revisit this. Otherwise, don't over-engineer.

The meta-lesson, one more time

Subagents are genuinely useful for a specific shape of problem. They're also the shiniest thing in the toolkit right now. Asking *“does my problem have that shape?”* is the whole game.

What You Built Today

- ① **A cleaner workflow.** You're working in your own workspace. No more merge conflicts against the shared repo.
- ② **A framework for skills.** A routing rule (stable-about-you / within-a-paper / live) plus the honest view: skills are empirical, and most get abandoned.
- ③ **Two MCPs on your map.** Zotero (friction) and GitHub (capability). The two-reason rule to evaluate others.
- ④ **An honest view of subagents.** Know they exist. Know the shape of problem they solve. Don't reach for them until your problem has that shape.

The thread running through

Agents solve symptoms well. Your job is to notice when the structure is wrong, when the skill is ceremony, when the MCP is novelty, and when the subagent is overkill.

Session 5: Expert Judgment and Capstone.

- Failure patterns — what goes wrong and how to recognize it
- When to correct the approach, restart the conversation, or do it yourself
- Capstone: build an end-to-end research artifact using everything from Sessions 1–4

Before next session:

- ➊ **Point your Zotero MCP at your own library** and run today's lit-review prompt on a real topic from your research
- ➋ **Run the routing test** on one thing you wished Claude already knew — is it a skill, project context, or MCP?
- ➌ **Try one thing from today** on your actual research: a skill, an MCP, or just the new repo workflow

Bring a question, a frustration, or a small win from the week. Session 5 opens with those.

